ENI → ENA → EFA

- Upto 10 GBPS
- VMDq
- TCP/IP
- Multiple ENI/instance
- Traffic can traverse across subnets
- VPC Networking, General purpose
- Default

- Upto 25 GBPS
- SR-IOV
- TCP/IP
- Single setting/per instance
- Traffic can traverses across subnets
- Low latency apps
- Optional on supported instance type

- Upto 100 GBPS
- OS-Bypass
- SRD
- One EFA per instance
- OS Bypass traffic is limited to single subnet and is not routable
- HPC and ML Apps
- Optional on supported instance type

Exam tip: HPC & Machine Learning >>>>>> EFA

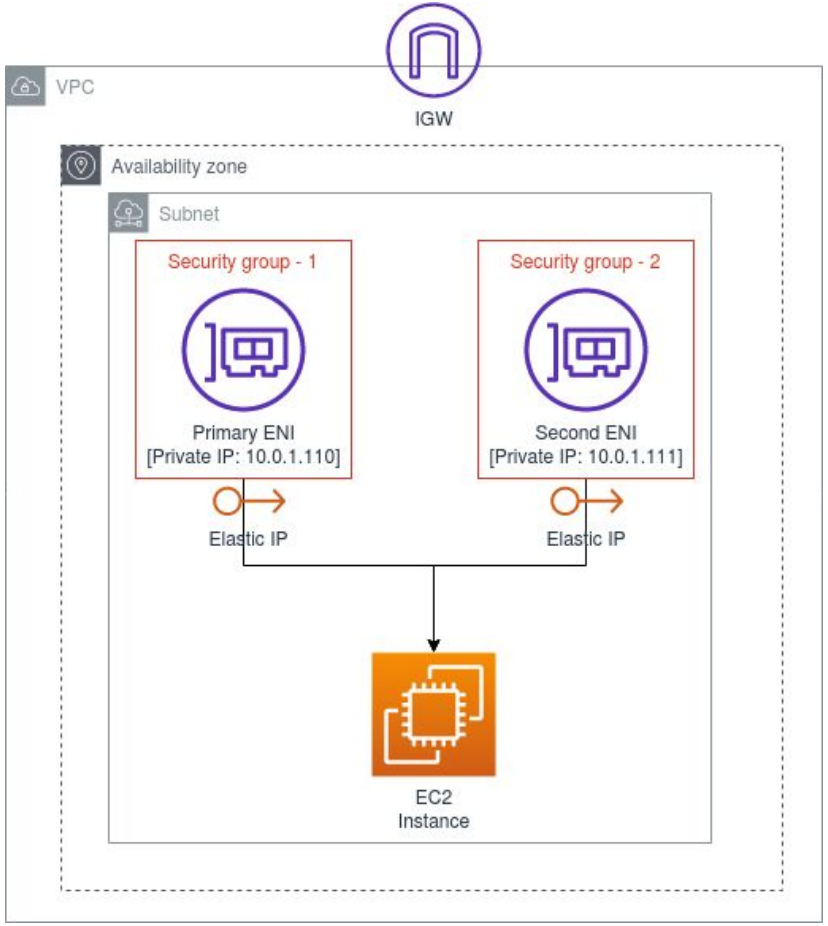

What will be migrated ?

Which Service wiil be used ?

On-Premises

DATA → DataSync

DATABASE → Database Migration Service

VM SERVER → Server Migration Service

Physical data transfer device

Snowball

Migration Hub

Real time Streaming

Capture

Transfer/Load

Analyze

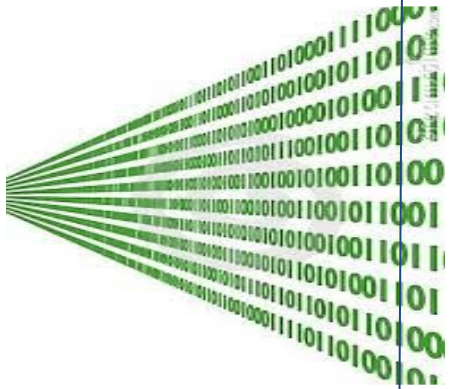

Kinesis Video Stream

Kinesis Streams

Kinesis Data Stream

Kinesis Firehose

Load streaming data into Amazon S3, Amazon Redshift, and Amazon Elasticsearch Service

Kinesis Analytics

Analyze data streams using standard SQL queries

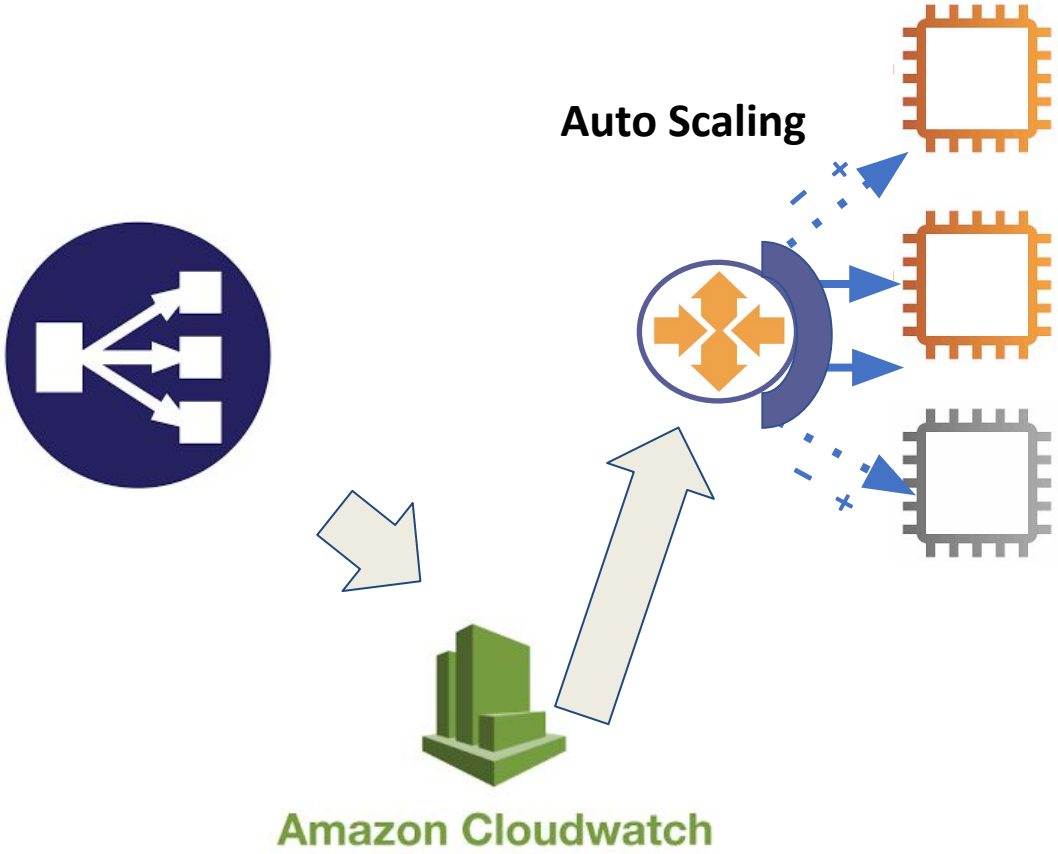

CPU utilization

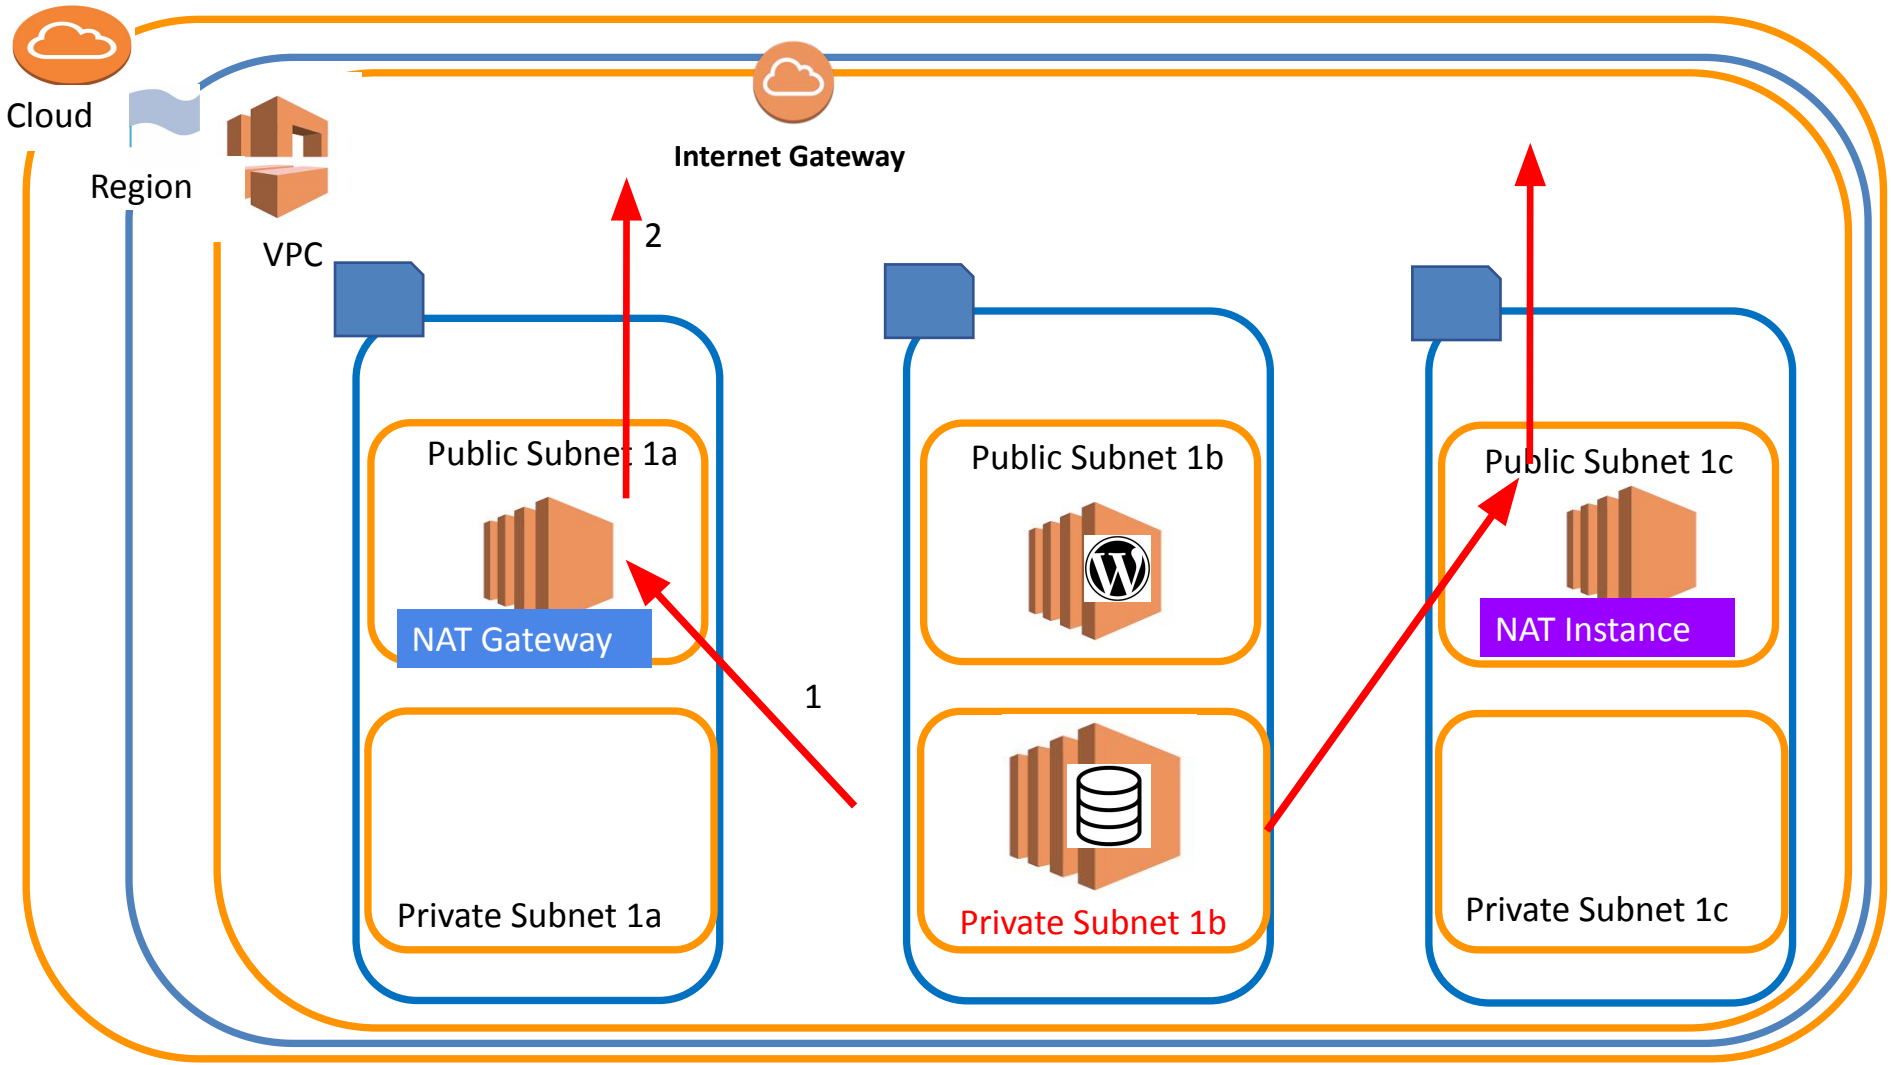
Cloud
Region
VPC
Internet Gateway
2
1
Public Subnet 1a
NAT Gateway
Private Subnet 1a
Public Subnet 1b
Private Subnet 1b
Public Subnet 1c
NAT Instance
Private Subnet 1c

Data Center

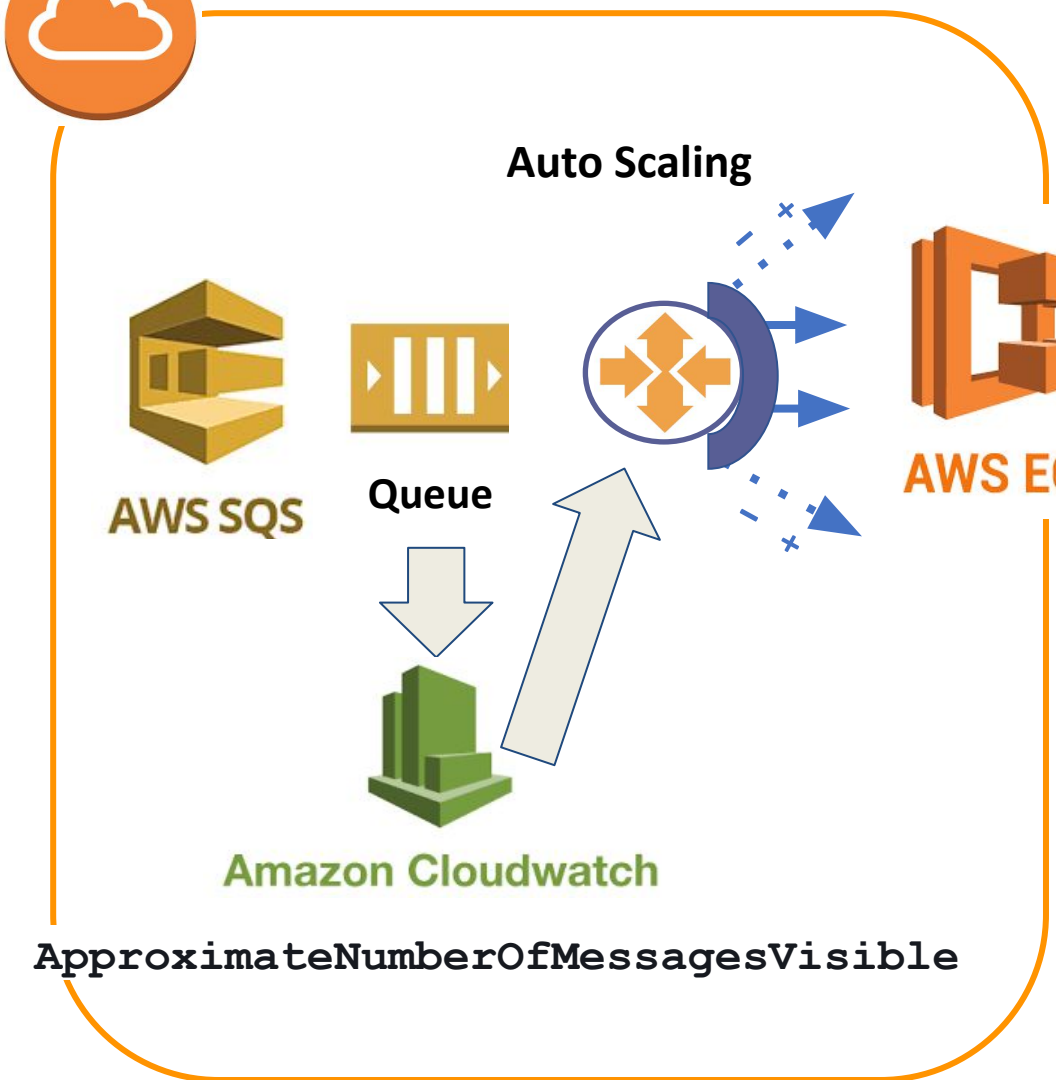

`ApproximateNumberOfMessagesVisible`

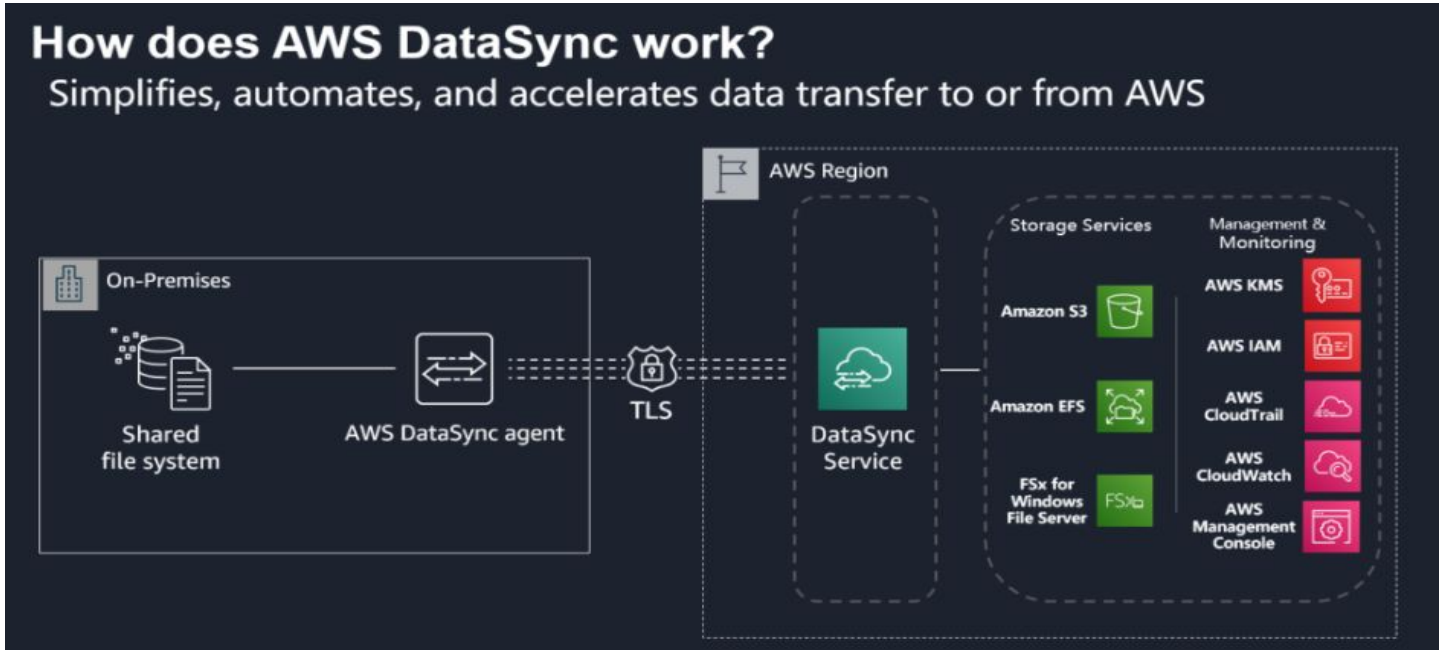

- **Public endpoints**:
- **Federal Information Processing Standard (FIPS) endpoints**: USA and Canada
- **Virtual private cloud (VPC) endpoints**: If you use a VPC endpoint, all communication from DataSync to AWS occurs through the endpoint in your AWS VPC. This establishes a private connection between your self-managed storage system, your VPC, and AWS services, providing extra security as your data is copied over the network.

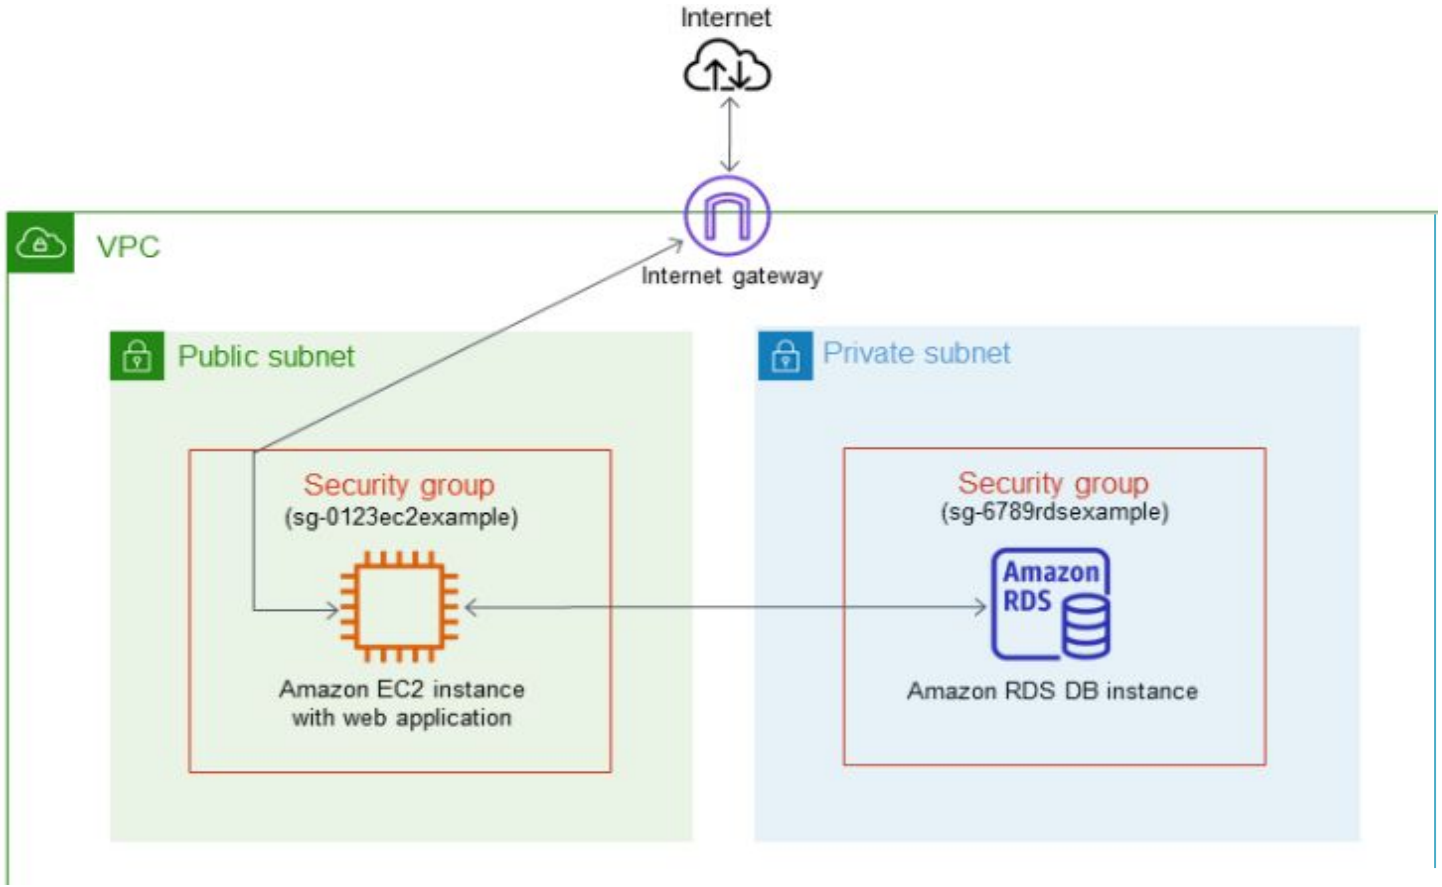
Internet
VPC
Internet gateway
Public subnet
Private subnet
Security group
(sg-0123ec2example)
Amazon EC2 instance
with web application
Security group
(sg-6789rdsexample)
Amazon
RDS
Amazon RDS DB instance

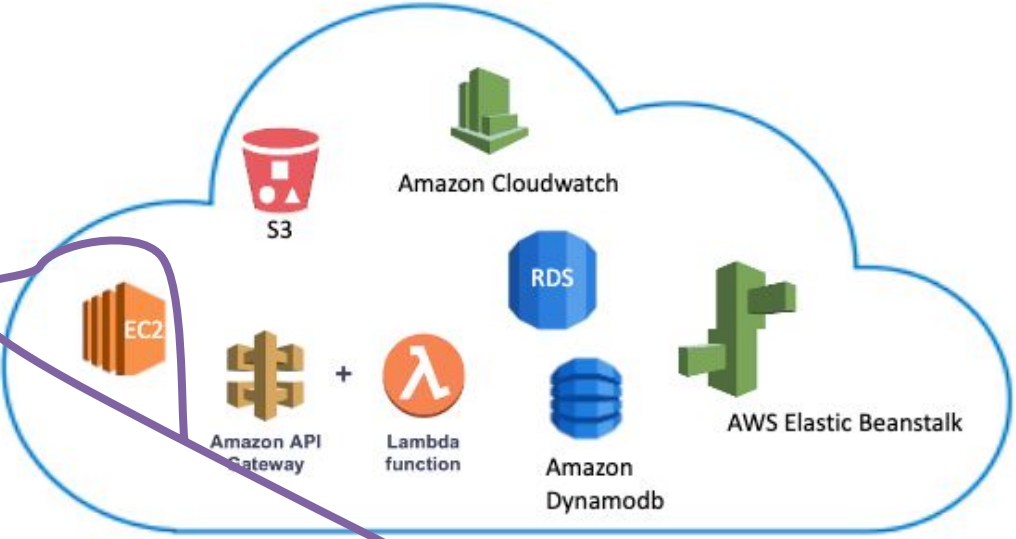
Amazon Cloudwatch
S3
RDS
EC2
+
Amazon API Gateway
Lambda function
Amazon Dynamodb
AWS Elastic Beanstalk


aws

IAM Role

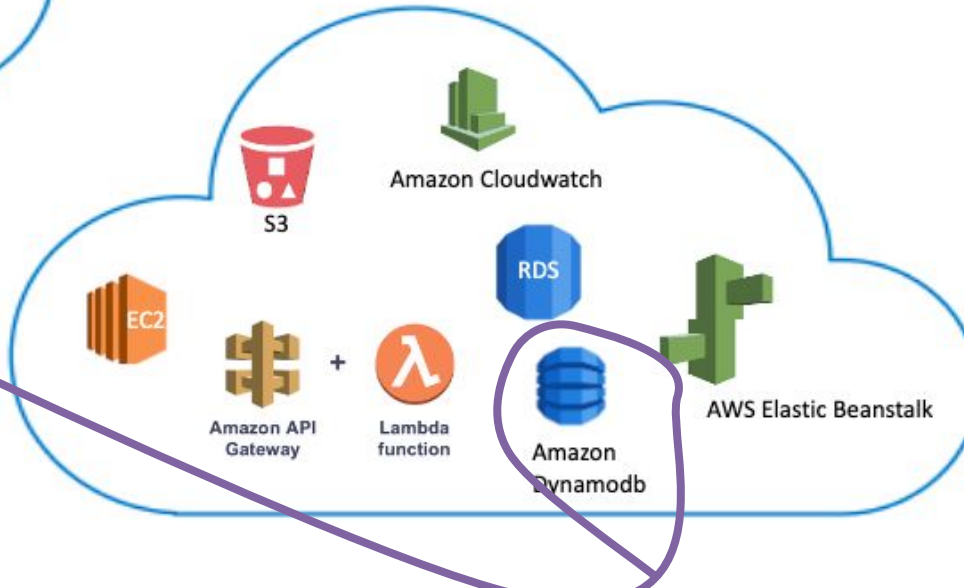
Amazon Cloudwatch
S3
RDS
EC2
+
Amazon API Gateway
Lambda function
Amazon Dynamodb
AWS Elastic Beanstalk

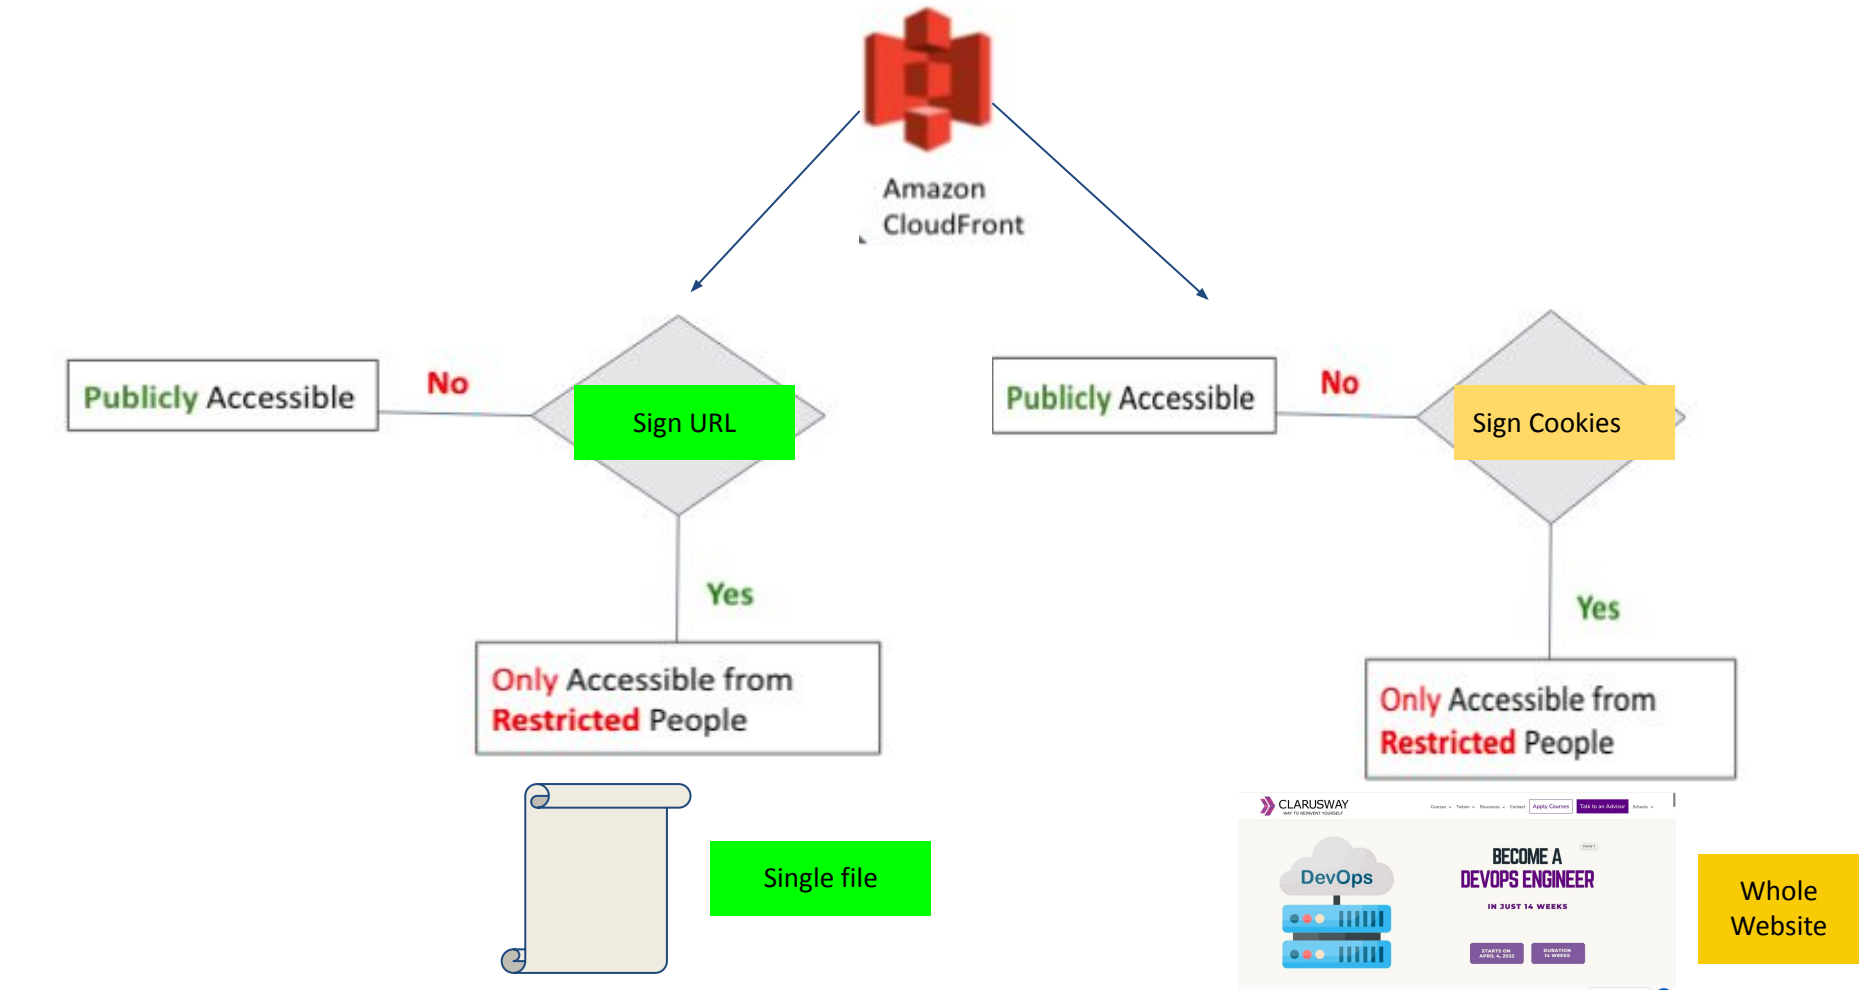
Amazon
CloudFront
Publicly Accessible
No
Sign URL
Yes
Only Accessible from
Restricted People
Single file
Publicly Accessible
No
Sign Cookies
Yes
Only Accessible from
Restricted People
CLARUSWAY
DevOps
BECOME A
DEVOPS ENGINEER
IN JUST 14 WEEKS
Whole
Website

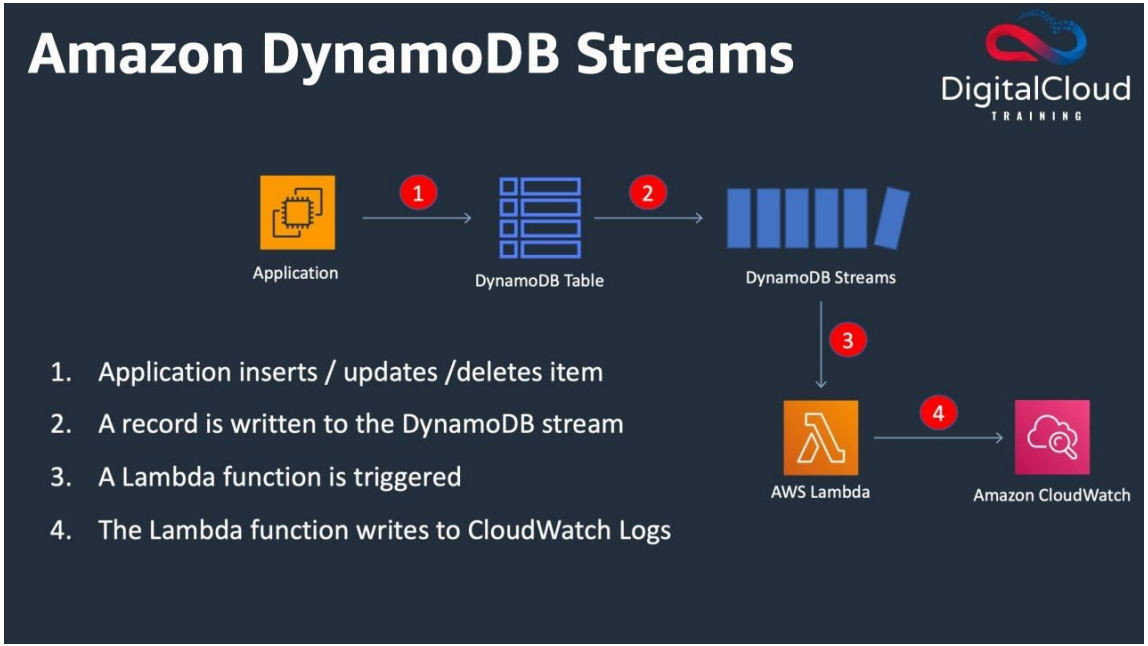

When enabled, DynamoDB Streams captures a time-ordered sequence of item-level modifications in a DynamoDB table and durably stores the information for **up to 24 hours. Applications can access a series of stream *records*, which contain an item change, from a DynamoDB stream in near real time.**

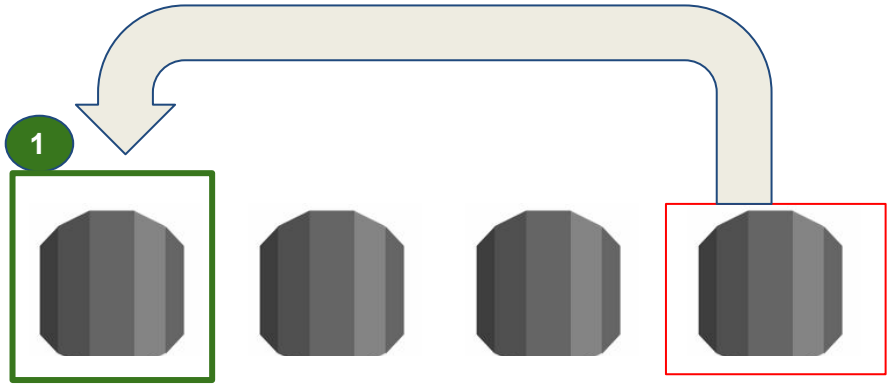

Point in time recovery

DESIRED :

23 JULY 2021
08:30:45

Snapshot

Transaction Logs

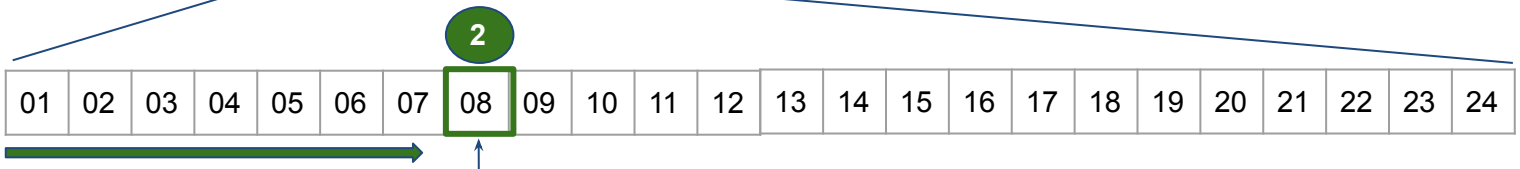

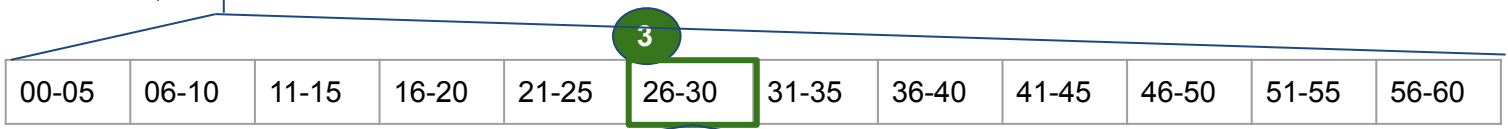

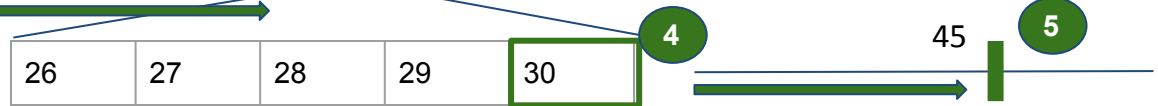

S3 static Web Hosting

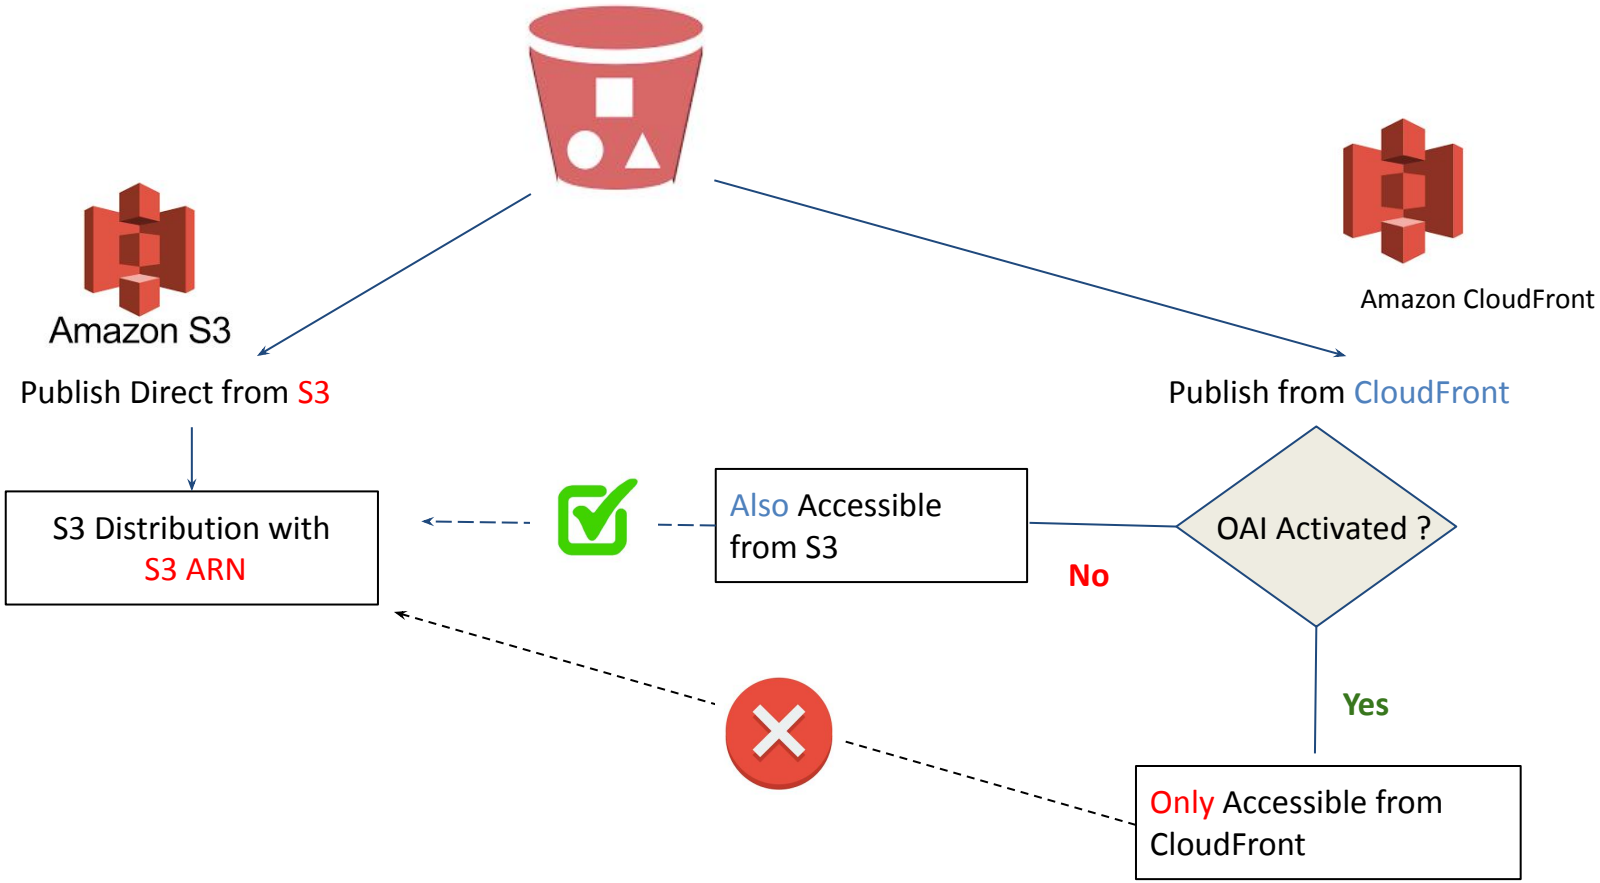

Real time Streaming

Capture

Transfer/Load

Analyze

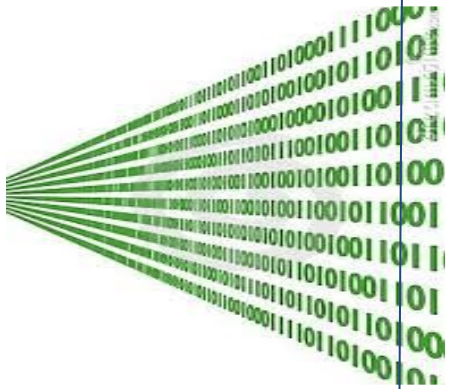

Kinesis Video Stream

Kinesis Streams

Kinesis Data Stream

Kinesis Firehose

Load streaming data into Amazon S3, Amazon Redshift, and Amazon Elasticsearch Service

Kinesis Analytics

Analyze data streams using standard SQL queries

What will be migrated ?

Which Service wiil be used ?

On-Premises

DATA → DataSync

DATABASE → Database Migration Service

VM SERVER → Server Migration Service

Physical data transfer device

Snowball

Migration Hub

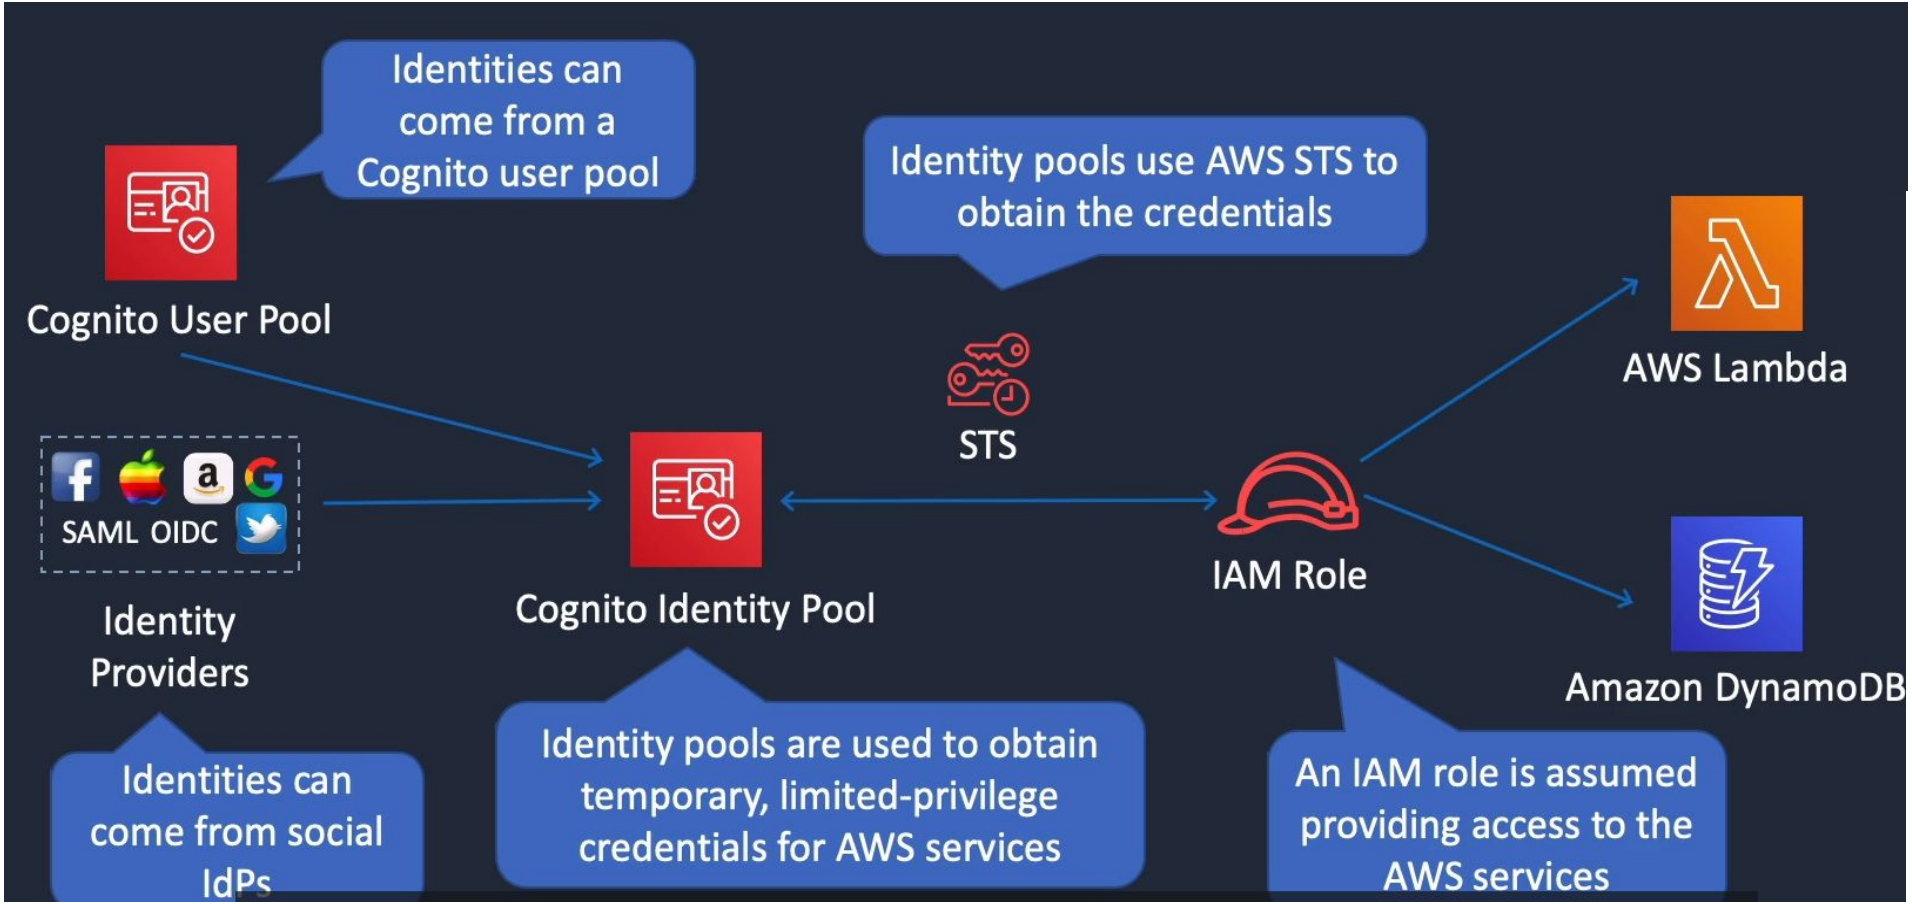
Identities can come from a Cognito user pool
Cognito User Pool
Identity pools use AWS STS to obtain the credentials
STS
AWS Lambda
SAML OIDC
Identity Providers
Cognito Identity Pool
IAM Role
Amazon DynamoDB
Identities can come from social IdPs
Identity pools are used to obtain temporary, limited-privilege credentials for AWS services
An IAM role is assumed providing access to the AWS services

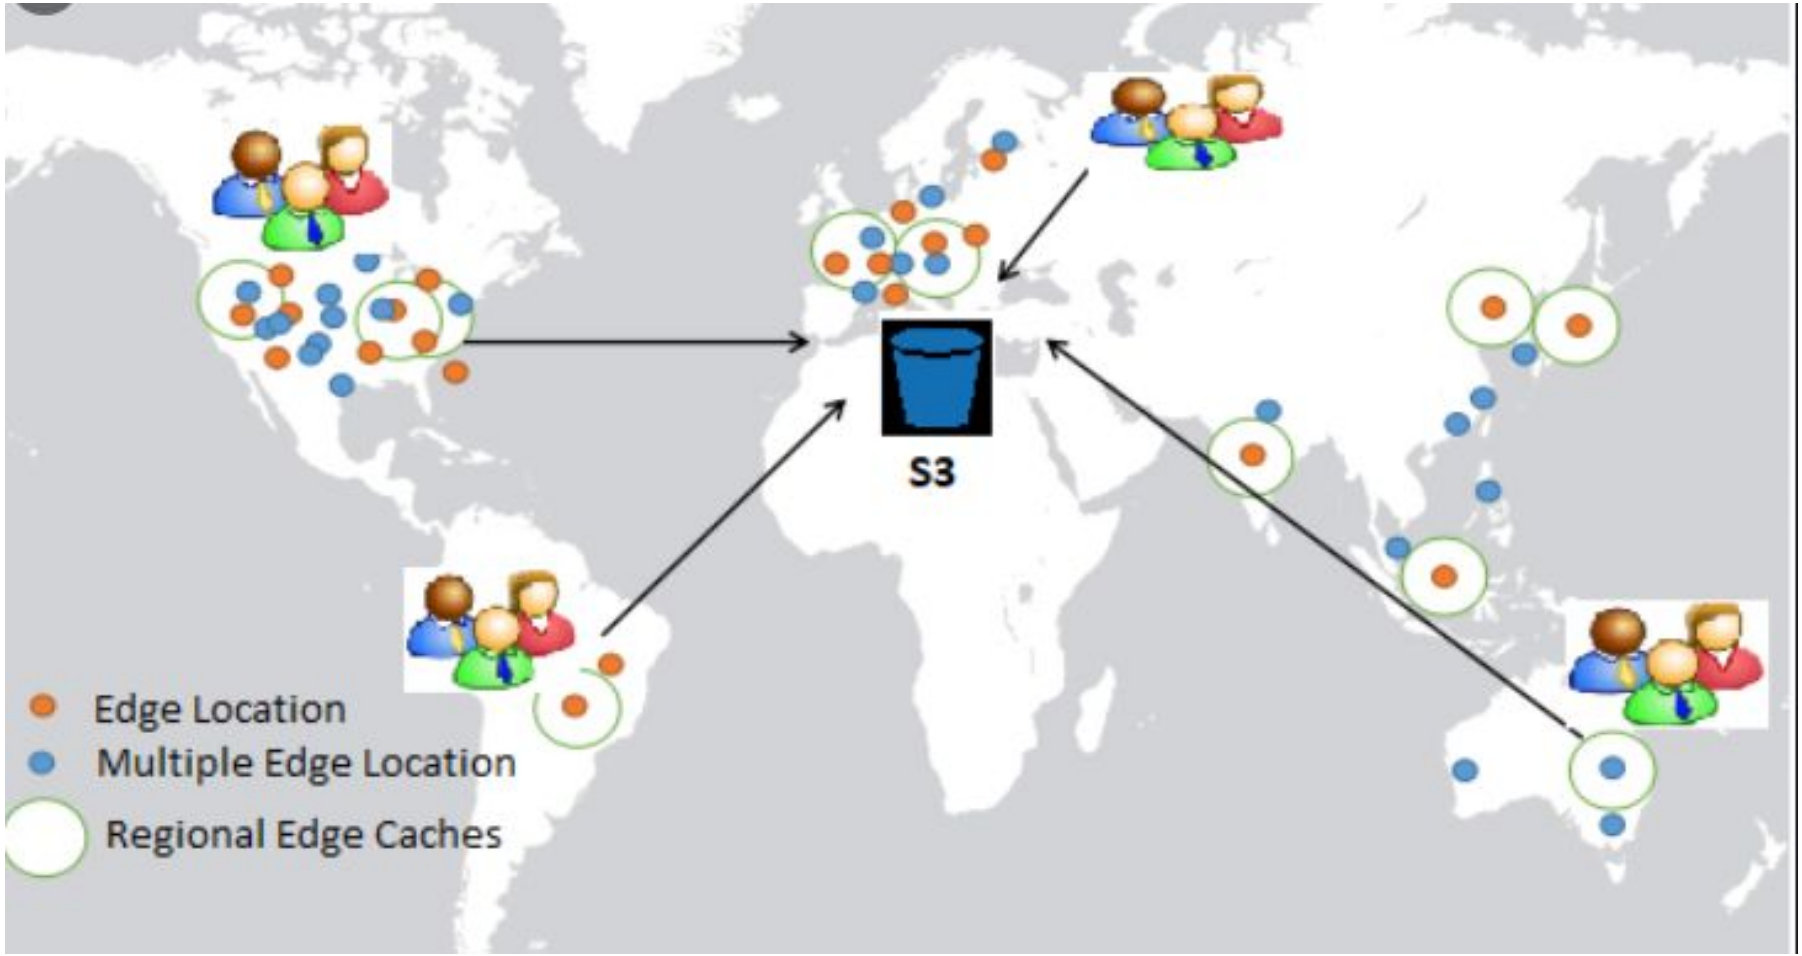
S3
Edge Location
Multiple Edge Location
Regional Edge Caches

If you have resources in multiple Availability Zones and they share one NAT Gateway, in the event that the NAT Gateway's Availability Zone is down, resources in the other Availability Zones lose internet access. To create an Availability Zone-independent architecture, create a NAT Gateway in each Availability Zone and configure your routing to ensure that resources use the NAT Gateway in the same Availability Zone.

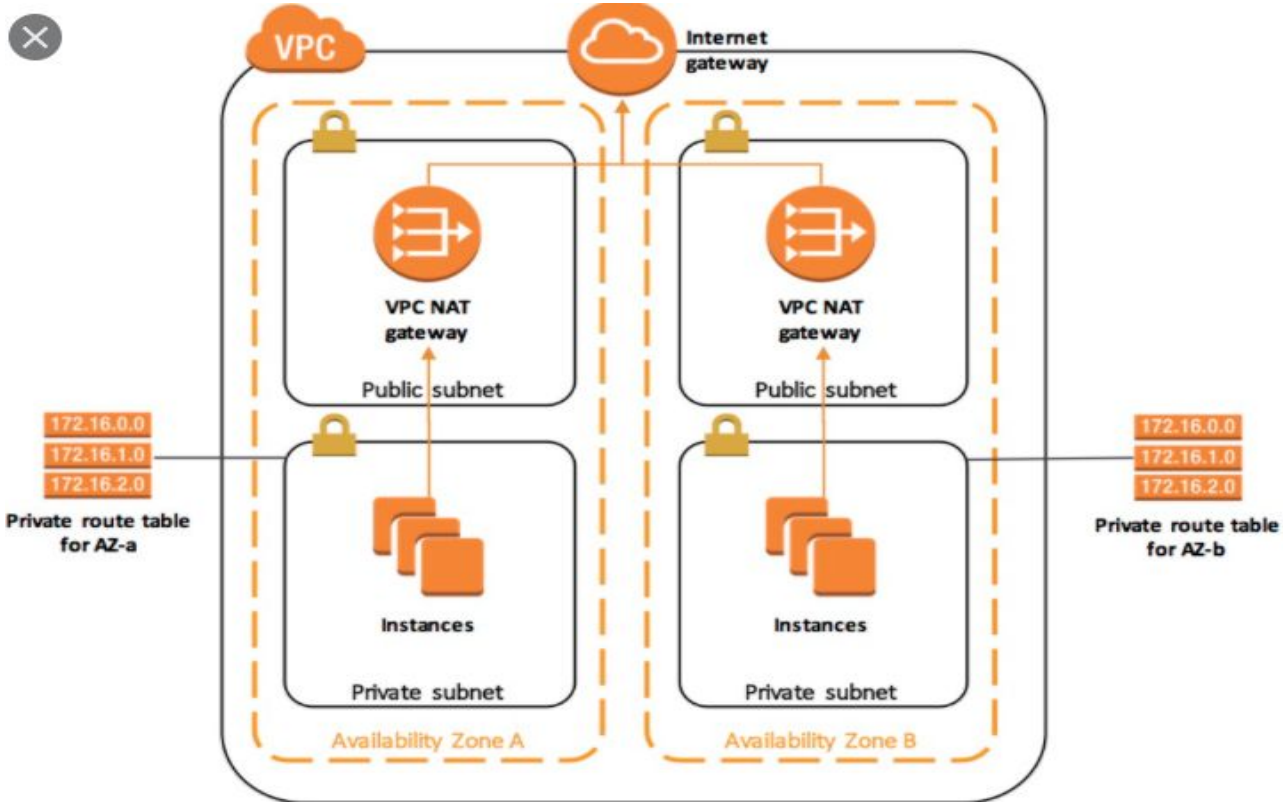

Amazon Aurora Key Features

Aurora Feature	Benefit
High performance and scalability	Offers high performance, self-healing storage that scales up to 64TB, point-in-time recovery and continuous backup to S3
DB compatibility	Compatible with existing MySQL and PostgreSQL open source databases
Aurora Replicas	In-region read scaling and failover target – up to 15 (can use Auto Scaling)
MySQL Read Replicas	Cross-region cluster with read scaling and failover target – up to 5 (each can have up to 15 Aurora Replicas)
Global Database	Cross-region cluster with read scaling (fast replication / low latency reads). Can remove secondary and promote
Multi-Master	Scales out writes within a region. In preview currently and will not appear on the exam
Serverless	On-demand, autoscaling configuration for Amazon Aurora - does not support read replicas or public IPs (can only access through VPC or Direct Connect - not VPN)

Throughput **5x MySQL RDS**
3x PostgreSQL RDS

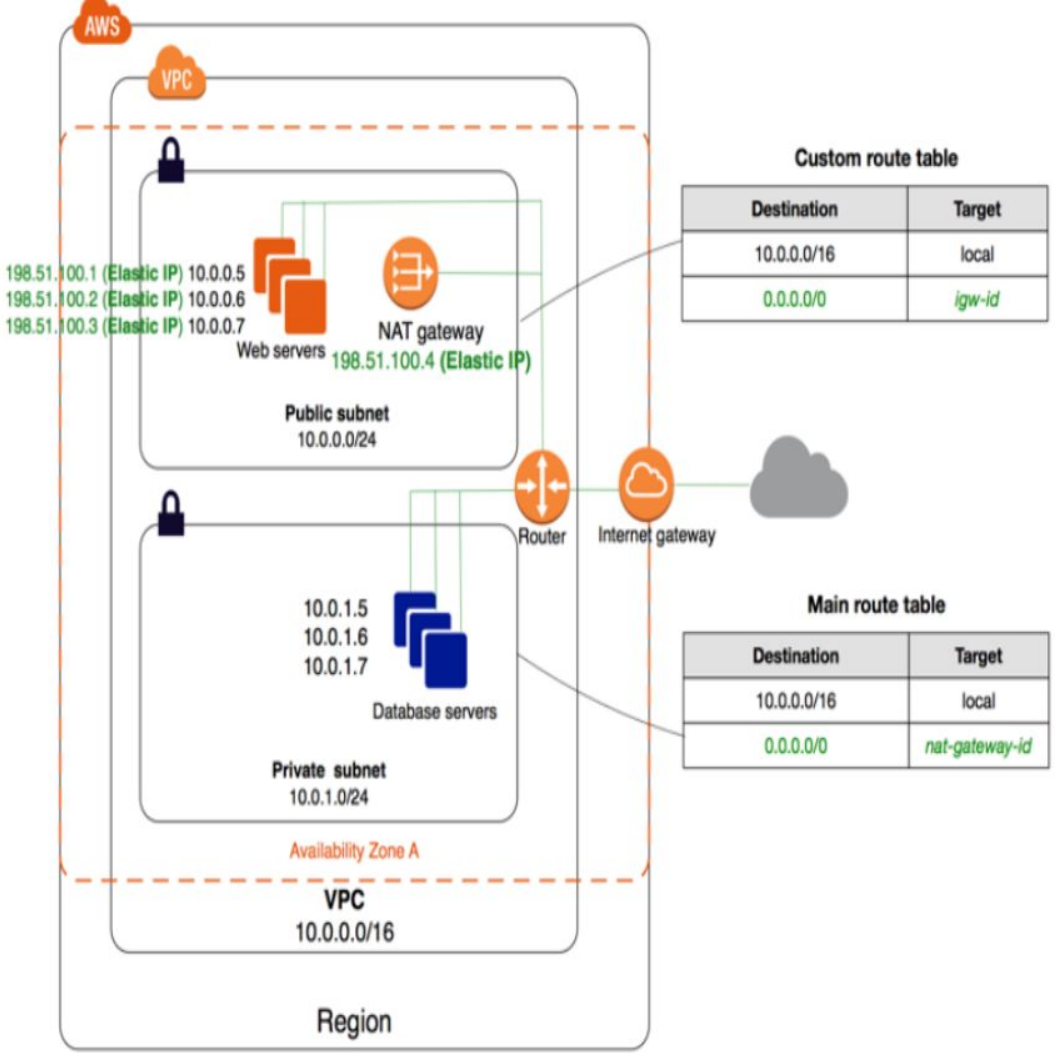
AWS
VPC
198.51.100.1 (Elastic IP) 10.0.0.5
198.51.100.2 (Elastic IP) 10.0.0.6
198.51.100.3 (Elastic IP) 10.0.0.7
Web servers
NAT gateway
198.51.100.4 (Elastic IP)
Public subnet
10.0.0.0/24
Router
Internet gateway
10.0.1.5
10.0.1.6
10.0.1.7
Database servers
Private subnet
10.0.1.0/24
Availability Zone A
VPC
10.0.0.0/16
Region
Custom route table
Destination
Target
10.0.0.0/16
local
0.0.0.0/0
igw-id
Main route table
Destination
Target
10.0.0.0/16
local
0.0.0.0/0
nat-gateway-id

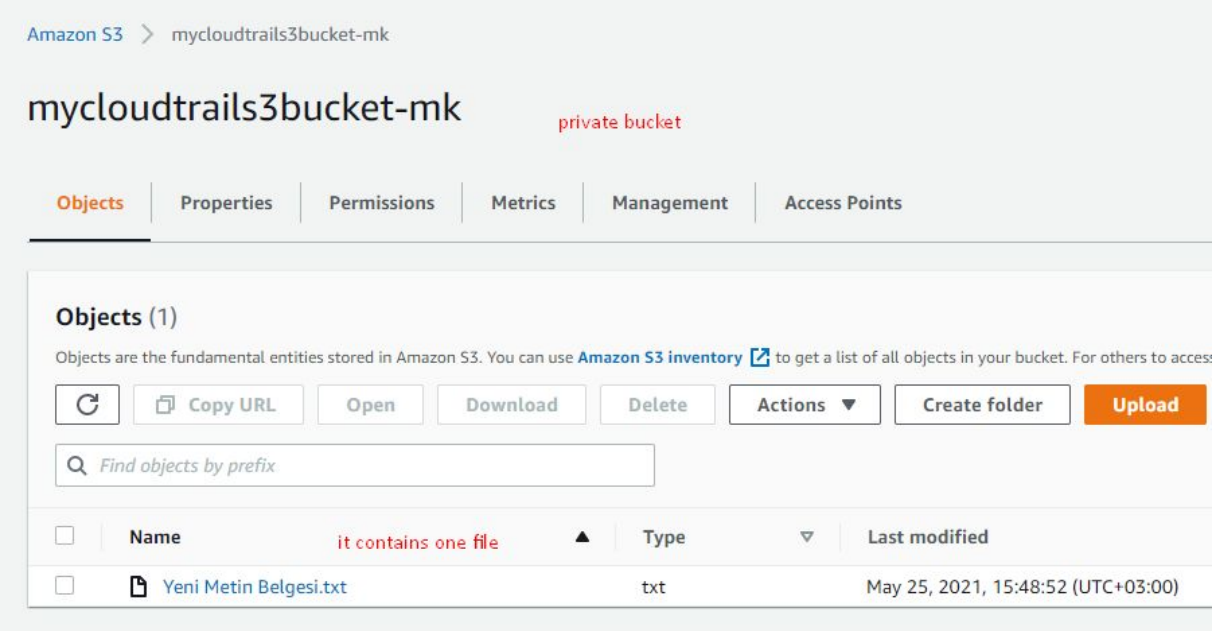


```
$ aws s3 presign s3://mycloudtrails3bucket-mk/"Yeni Metin Belgesi.txt" --expires-in 30
https://mycloudtrails3bucket-mk.s3.us-east-1.amazonaws.com/Yeni%20Metin%20Belgesi.txt?X-Amz-
Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAY3YMFVJJYLFDLWV5%2F20210525%2Fus-east-1%2Fs3
%2Faws4_request&X-Amz-Date=20210525T125743Z&X-Amz-Expires=30&X-Amz-SignedHeaders=host&X-Amz-
Signature=b02d2c978829ff3bed34001380b8b9e0b51262727d2b37391c4b5c12e0d6f31c
```

In the question : there will be a rds mysql and is expected a high number of read and writes. Which ebs types supports this?

	Solid State Drives (SSD)					Hard Disk Drives (HDD)	
Volume Type	EBS Provisioned IOPS SSD (io2 Block Express)	EBS Provisioned IOPS SSD (io2)	EBS Provisioned IOPS SSD (io1)	EBS General Purpose SSD (gp3) announced Dec 1, 2020	EBS General Purpose SSD (gp2)*	Throughput Optimized HDD (st1)	Cold HDD (sc1)
Short Description	Highest performance SSD volume designed for business-critical latency-sensitive **transactional workloads**	**Highest performance** and highest durability SSD volume designed for **latency-sensitive transactional workloads**	Highest performance SSD volume designed for **latency-sensitive transactional workloads**	**Lowest cost SSD volume that balances price performance for a wide variety of transactional workloads**	General Purpose SSD volume that **balances price performance for a wide variety of transactional workloads**	**Low cost** HDD volume designed for **frequently** accessed, **throughput intensive workloads**	**Lowest cost** HDD volume designed for **less frequently accessed workloads**
Durability	99.999%		99.8% - 99.9% durability			99.8% - 99.9% durability	
Use Cases	Largest, most I/O intensive, mission critical deployments of NoSQL and relational **databases** such as Oracle, SAP HANA, Microsoft SQL Server, and SAS Analytics	I/O-intensive NoSQL and relational **databases**	I/O-intensive NoSQL and relational **databases**	Virtual desktops, medium sized single instance **databases** such as Microsoft SQL Server and Oracle, latency sensitive interactive applications, boot volumes, and dev/test environments	Virtual desktops, medium sized single instance **databases** such as Microsoft SQL Server and Oracle, latency sensitive interactive applications, boot volumes, and dev/test environments	**Big data, data warehouses, log processing**	**Colder data requiring fewer scans per day**

With Amazon RDS, you can create a MariaDB, MySQL, Oracle, or PostgreSQL read replica in a different AWS Region from the source DB instance. Creating a cross-Region read replica isn't supported for SQL Server on Amazon RDS.

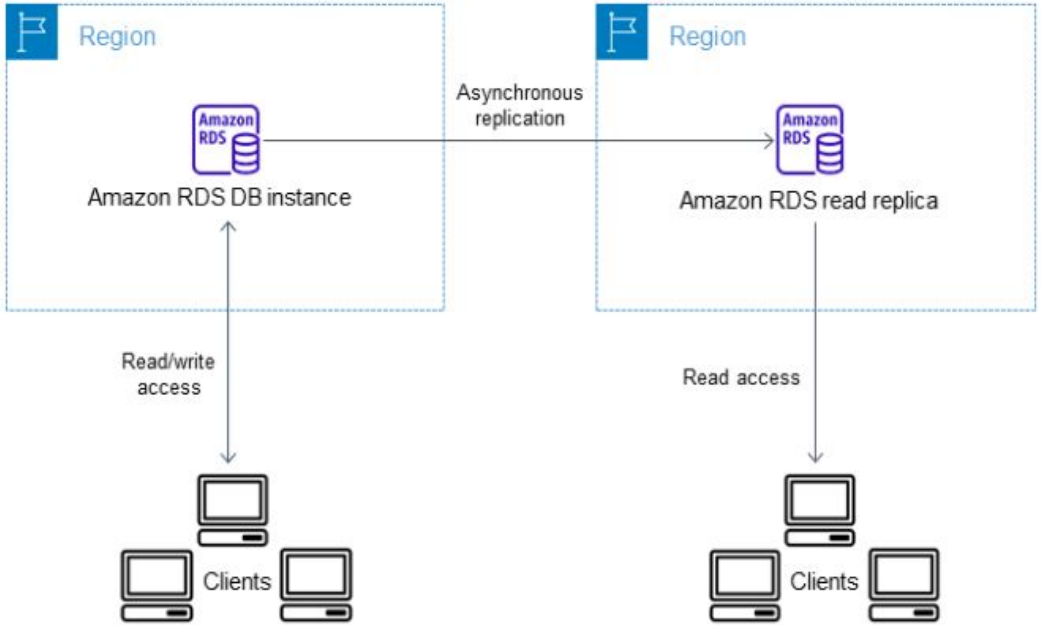

RDS Multi-AZ Failover

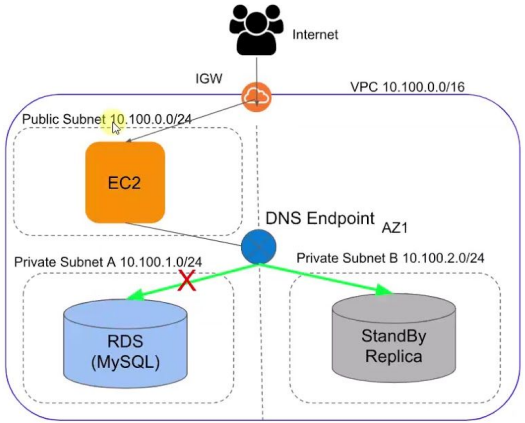

Availability - Failover-Senkron

Performance- Asenkron

Real time Streaming

Capture

Transfer/Load

Analyze

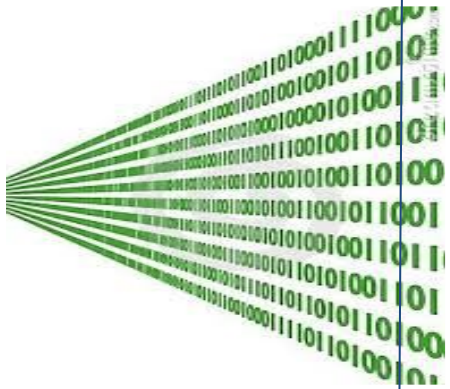

Kinesis Video Stream

Kinesis Streams

Kinesis Data Stream

Kinesis Firehose

Load streaming data into Amazon S3, Amazon Redshift, and Amazon Elasticsearch Service

Kinesis Analytics

Analyze data streams using standard SQL queries

1. Define Load Balancer | 2. Assign Security Groups | 3. Configure Security Settings | 4. Configure Health Check | **5. Add EC2 Instances** | 6. Add Tags | 7. Review

Step 5: Add EC2 Instances

The table below lists all your running EC2 Instances. Check the boxes in the Select column to add those instances to this load balancer.

VPC vpc-f52d178f (172.31.0.0/16) | default-vpc

Instance	Name	State	Security groups	Zone	Subnet ID	Subnet CIDR
i-0667836734e141b24	Cloudwatch_instance	stopped	launch-wizard-40	us-east-1d	subnet-ed49bccc	172.31.80.0/20
i-09a81ac9ed47f33a3	pet-clinic-demo-server	stopped	launch-wizard-21	us-east-1f	subnet-3246e63c	172.31.64.0/20
i-0c8b559e55c99cfe1	Jenkins Server of tomy-jenkinsserver-p...	stopped	tomy-jenkinsserver-petclinic-JenkinsS...	us-east-1f	subnet-3246e63c	172.31.64.0/20
i-030841b79c0f01e67	tyler-nexus-server	stopped	Nexus-SG	us-east-1f	subnet-3246e63c	172.31.64.0/20
i-090276ecb3211442a	Jenkins Server of walter-petclinic	stopped	walter-petclinic-JenkinsServerSecurity...	us-east-1f	subnet-3246e63c	172.31.64.0/20

Availability Zone Distribution

1 instance in us-east-1d

Cross-zone load balancing distributes traffic evenly across all targets in the Availability Zones enabled for the load balancer.

☑ Enable Cross-Zone Load Balancing ⓘ

☑ Enable Connection Draining ⓘ 300 seconds

Exam tip: DynamoDb and Metadata

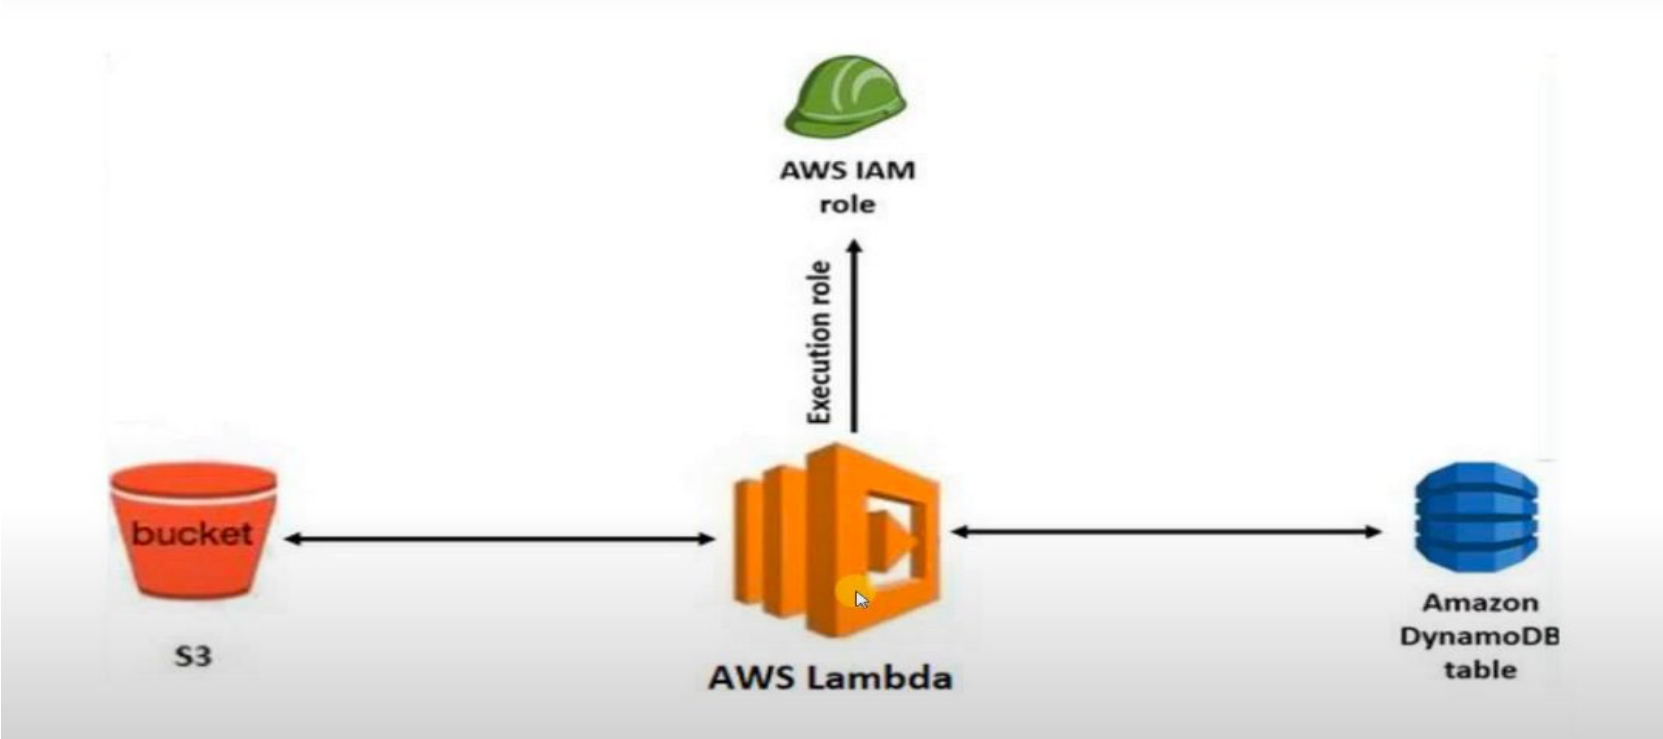

Introduction to Elastic Beanstalk

Why AWS Elastic Beanstalk?

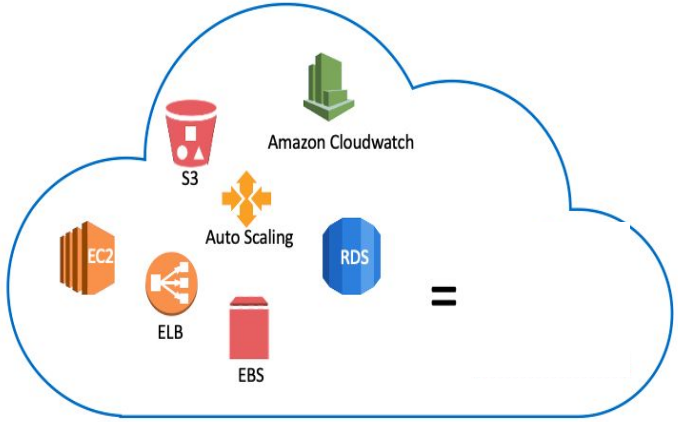

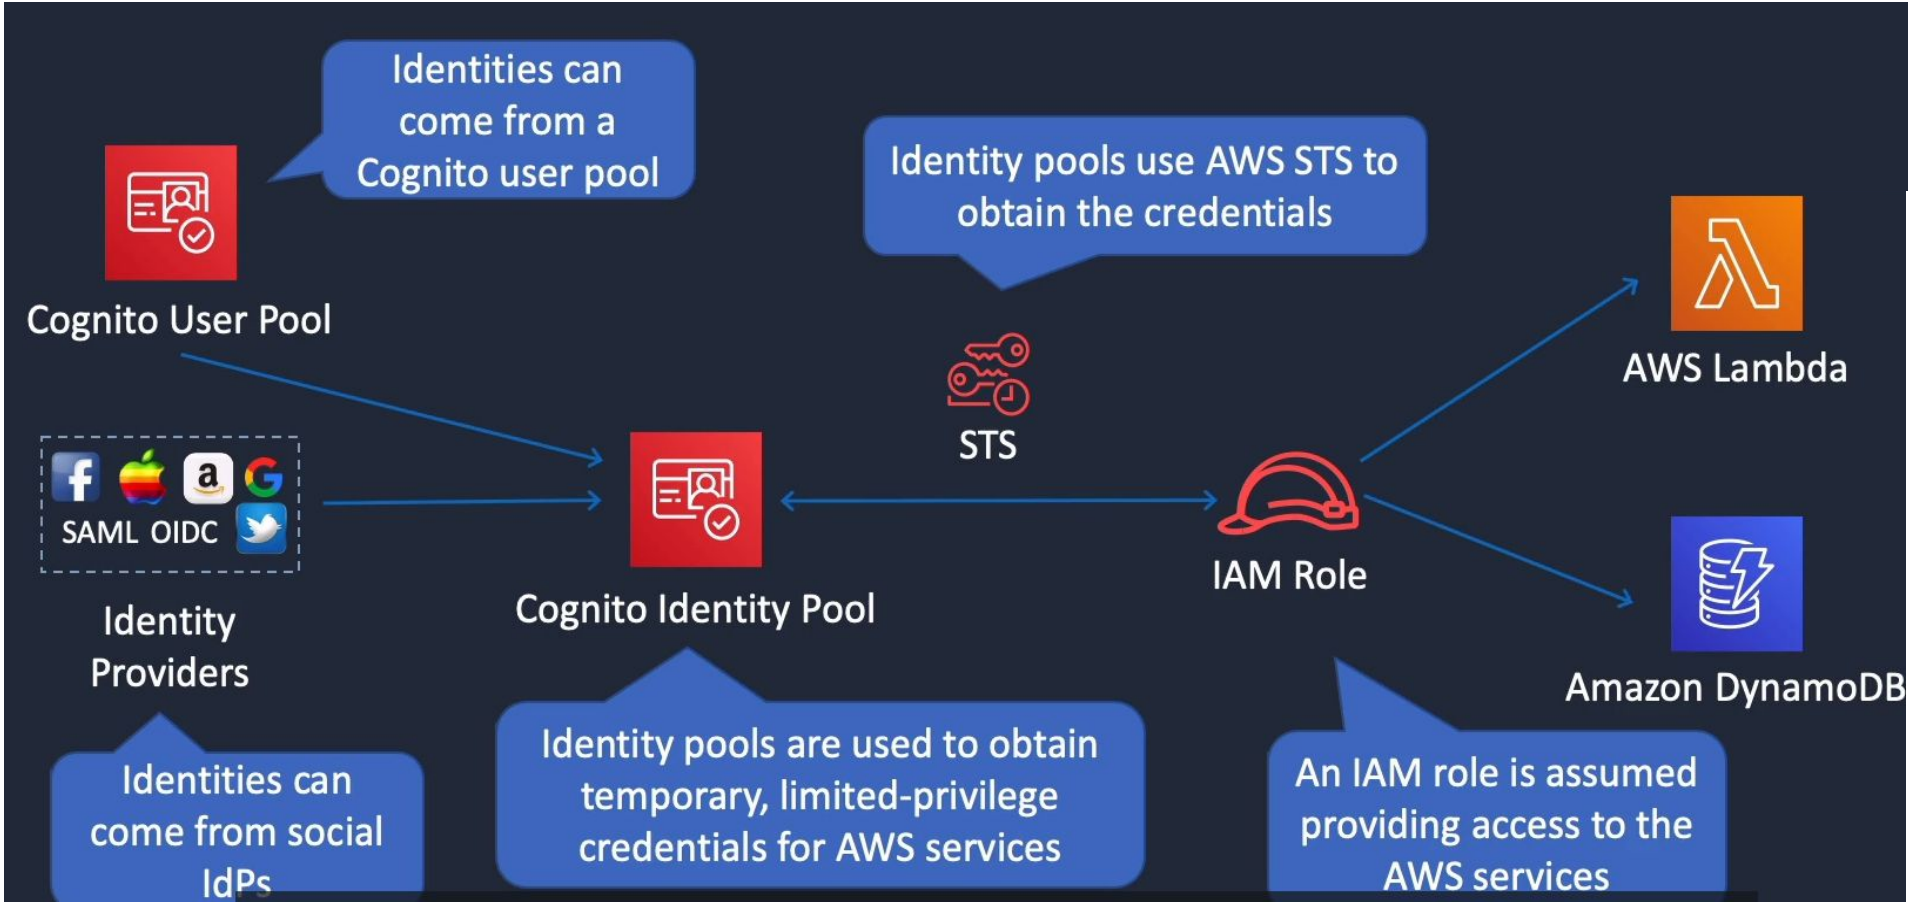
Identities can come from a Cognito user pool
Cognito User Pool
Identity pools use AWS STS to obtain the credentials
STS
AWS Lambda
SAML OIDC
Identity Providers
Cognito Identity Pool
IAM Role
Amazon DynamoDB
Identities can come from social IdPs
Identity pools are used to obtain temporary, limited-privilege credentials for AWS services
An IAM role is assumed providing access to the AWS services

Storage Classes

Standard IA (Infrequent Access)

Infrequently Accessed Data

- Standard IA (Infrequent Access) is a convenient for infrequently accessed files
- But in case of access, it provides you to reach the file quickly.
- In fact, it designed for the data which requires less frequent access, but with longer storage time than the Standard class
- It is cheaper than Standard class as long as you access **infrequently.**

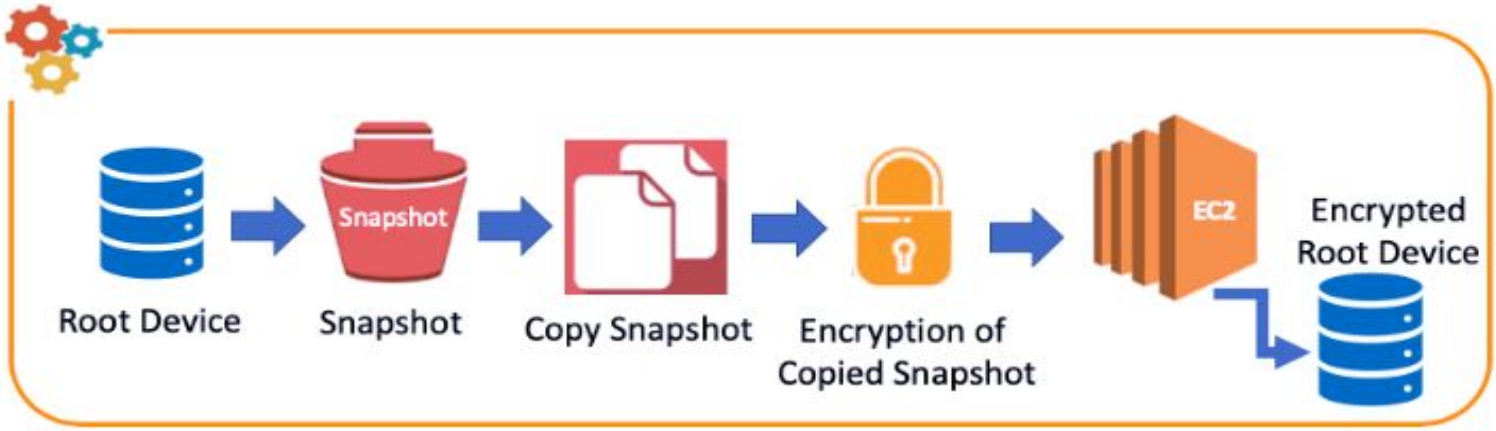
Snapshot
EC2
Encrypted
Root Device
Root Device
Snapshot
Copy Snapshot
Encryption of
Copied Snapshot

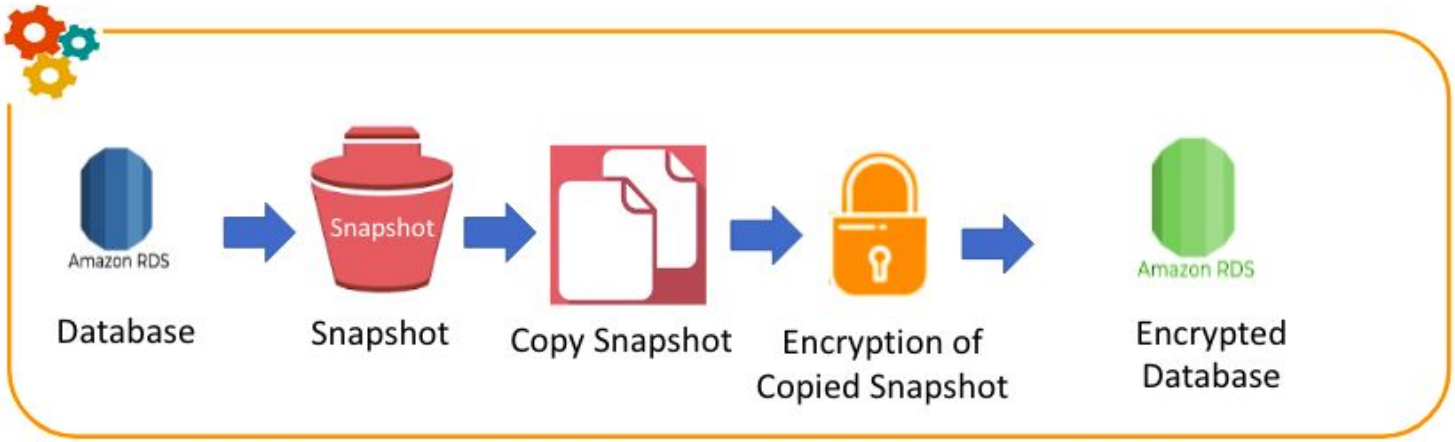
Amazon RDS
Snapshot
Amazon RDS
Database
Snapshot
Copy Snapshot
Encryption of
Copied Snapshot
Encrypted
Database

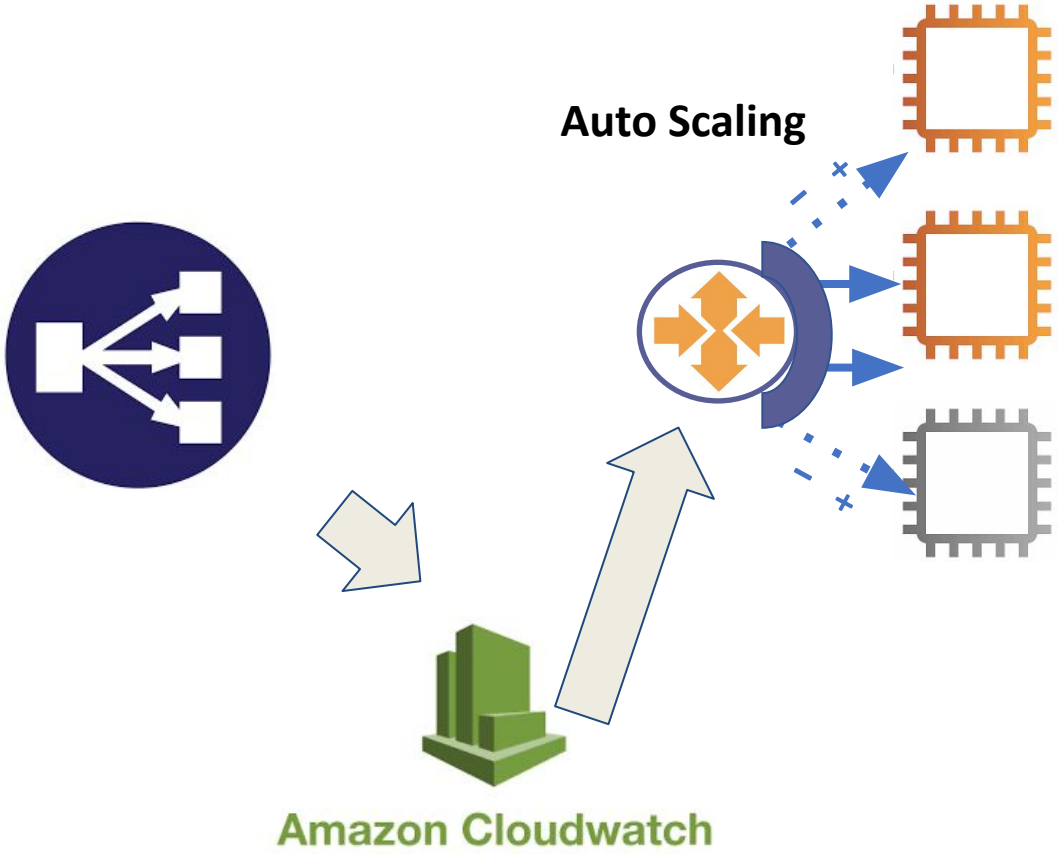

CPU utilization

MEMBER AND MASTER ACCOUNT LEAVING PROCESS

1. Remove the **member account** from the old Organization.
2. Send an invite to the member account from the new Organization.
3. Accept the invite to the new Organization from the member account.

---

4. Delete the old Organization.
5. Send an invite to the master account
6. Accept the invite to the new Organization from the master account

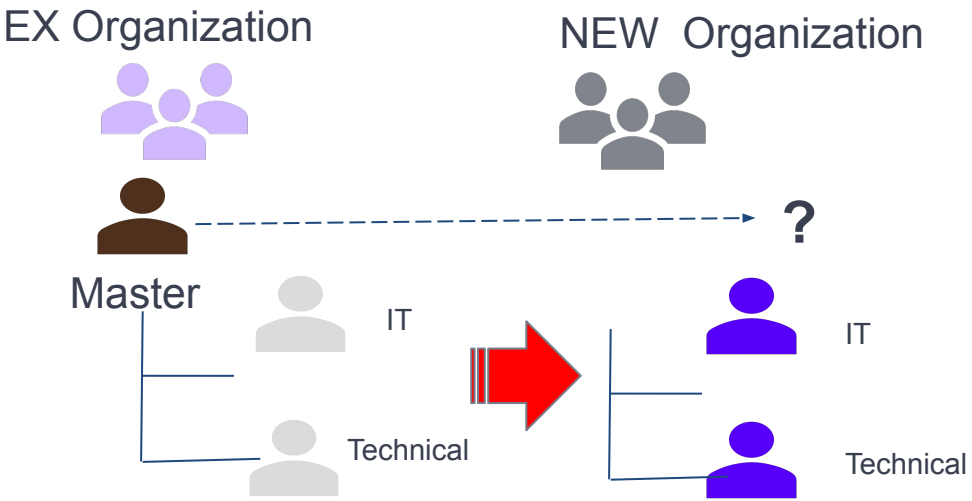

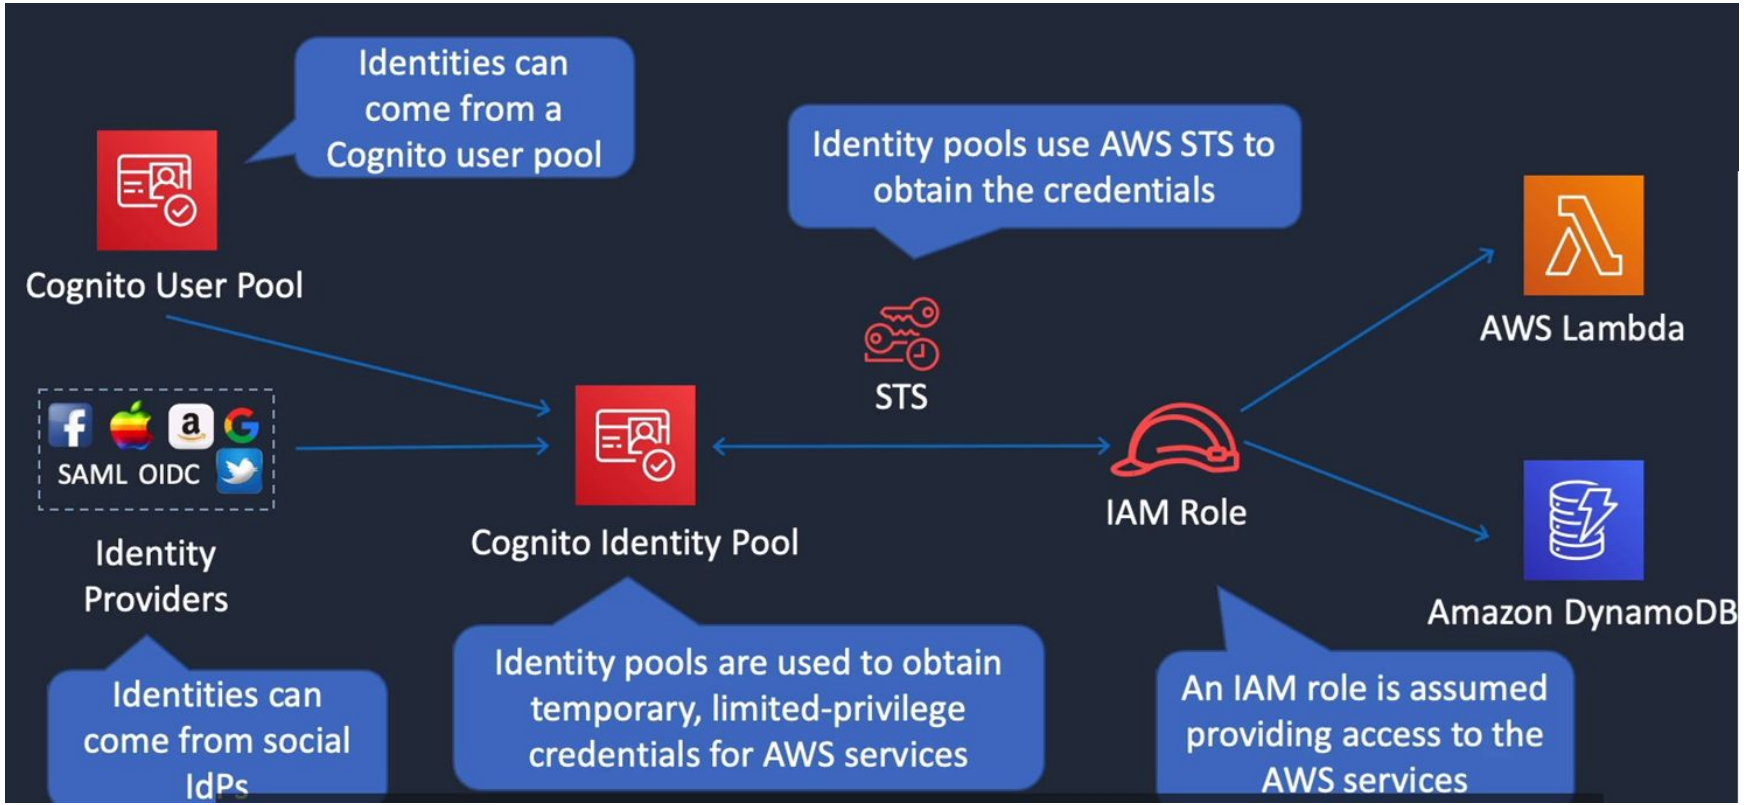
Identities can come from a Cognito user pool
Cognito User Pool
Identity pools use AWS STS to obtain the credentials
STS
AWS Lambda
SAML OIDC
Identity Providers
Cognito Identity Pool
IAM Role
Amazon DynamoDB
Identities can come from social IdPs
Identity pools are used to obtain temporary, limited-privilege credentials for AWS services
An IAM role is assumed providing access to the AWS services

Component	Description
Templates	The JSON or YAML text file that contains the instructions for building out the AWS environment
Stacks	The entire environment described by the template and created, updated, and deleted as a single unit
StackSets	AWS CloudFormation StackSets extends the functionality of stacks by enabling you to create, update, or delete stacks across multiple accounts and regions with a single operation
Change Sets	A summary of proposed changes to your stack that will allow you to see how those changes might impact your existing resources before implementing them

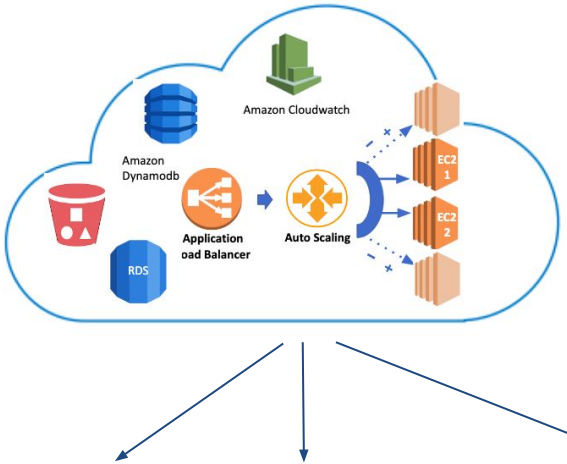

OpsWorks Stacks CloudFormation Elastic Beanstalk

Multi-Value Answer Policies let you configure Route53 to return multiple values such as IP addresses for your web-servers, in response to DNS queries.

Multiple values can be specified for almost any record. Route53 automatically performs health-checks on resources and only returns values of ones deemed healthy.

Similar to Simple Routing, however with an added health check for your record set resources.

Exam tip: Weighted Routing Policy is for blue-green deployment

Server Side Encryption (At Rest)

Encryption in Transition

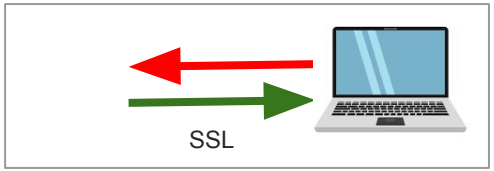

Client Side Encryption
(Transition+At Rest)

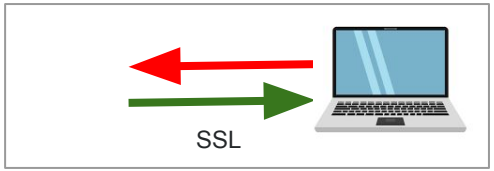

Failover Scenario

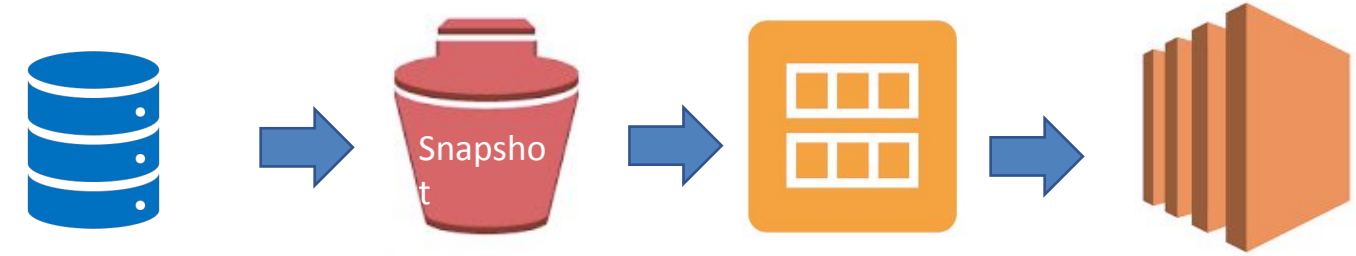

Additional configurations

Use defaults

These configurations are optional, and default settings have been defined to help you get started with your cluster. Turn off "Use defaults" to modify these settings now.

▼ Network and security

Virtual private cloud (VPC)
This VPC defines the virtual networking environment for this cluster. Choose a VPC that has a subnet group. Only valid VPCs are enabled in the list.

You can't change the VPC associated with this cluster after the cluster has been created. Learn more

VPC security groups
This VPC security group defines which subnets and IP ranges the cluster can use in the VPC.

Choose one or more security groups

default
sg-e82ceac9

Cluster subnet group
Choose the Amazon Redshift subnet group to launch the cluster in.

Availability Zone
Specify the Availability Zone that you want the cluster to be created in. Otherwise, Amazon Redshift chooses an Availability Zone for you.

No preference

Enhanced VPC routing
Forces cluster traffic through a VPC.

- (•) Disabled
- ( ) Enabled

Publicly accessible
Allow instances and devices outside the VPC connect to your database through the cluster endpoint.

- (•) No

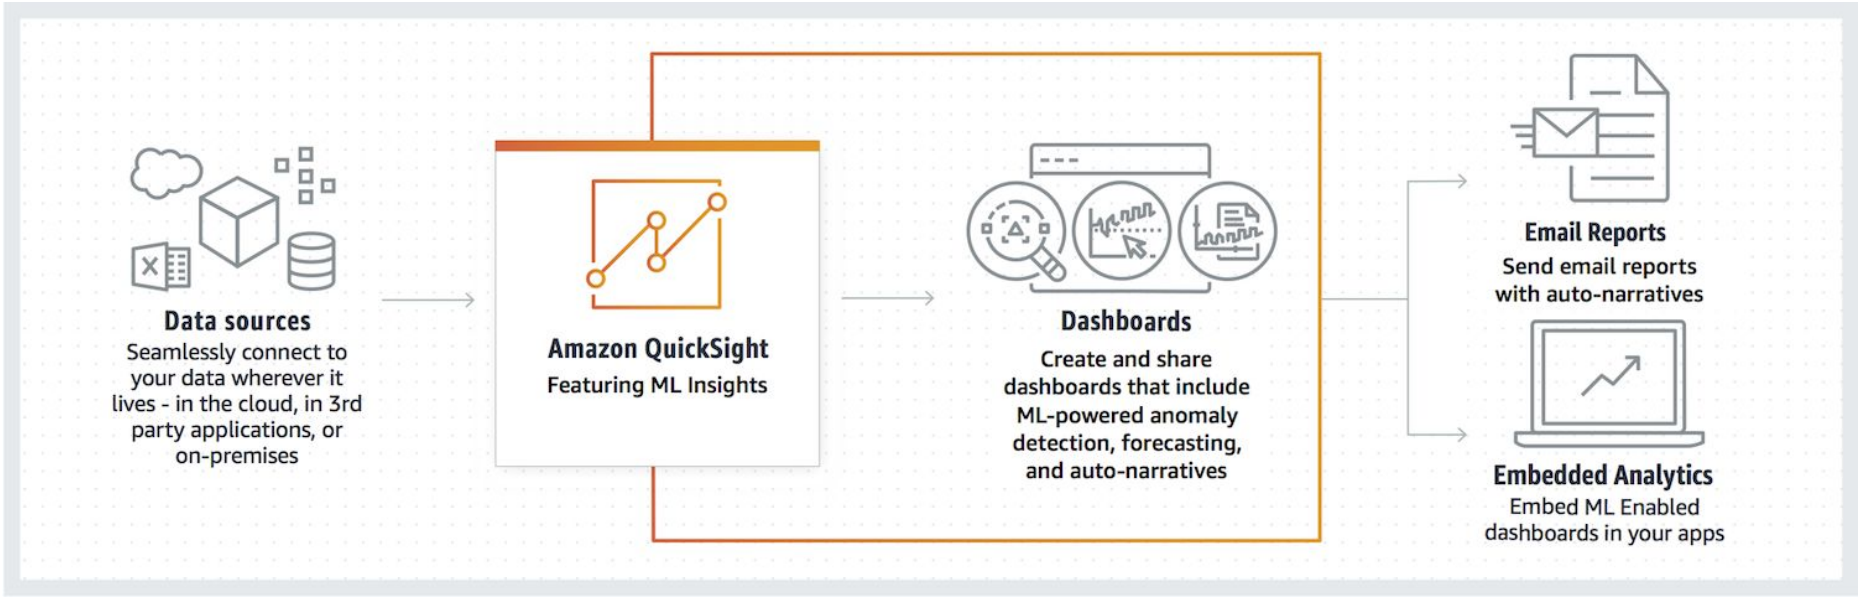
Data sources
Seamlessly connect to your data wherever it lives - in the cloud, in 3rd party applications, or on-premises
Amazon QuickSight
Featuring ML Insights
Dashboards
Create and share dashboards that include ML-powered anomaly detection, forecasting, and auto-narratives
Email Reports
Send email reports with auto-narratives
Embedded Analytics
Embed ML Enabled dashboards in your apps

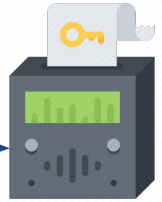

External

Custom key store
(CloudHSM)

Create function Info

Choose one of the following options to create your function.

Author from scratch
Start with a simple Hello World example.

Use a blueprint
Build a Lambda application from sample code and configuration presets for common use cases.

Container image
Select a container image to deploy for your function.

Browse serverless app repository
Deploy a sample Lambda application from the AWS Serverless Application Repository.

Basic information

Function name
Enter a name that describes the purpose of your function.

myFunctionName

Use only letters, numbers, hyphens, or underscores with no spaces.

Runtime Info
Choose the language to use to write your function. Note that the console code editor supports only Node.js, Python, and Ruby.

.NET Core 3.1 (C#/PowerShell)

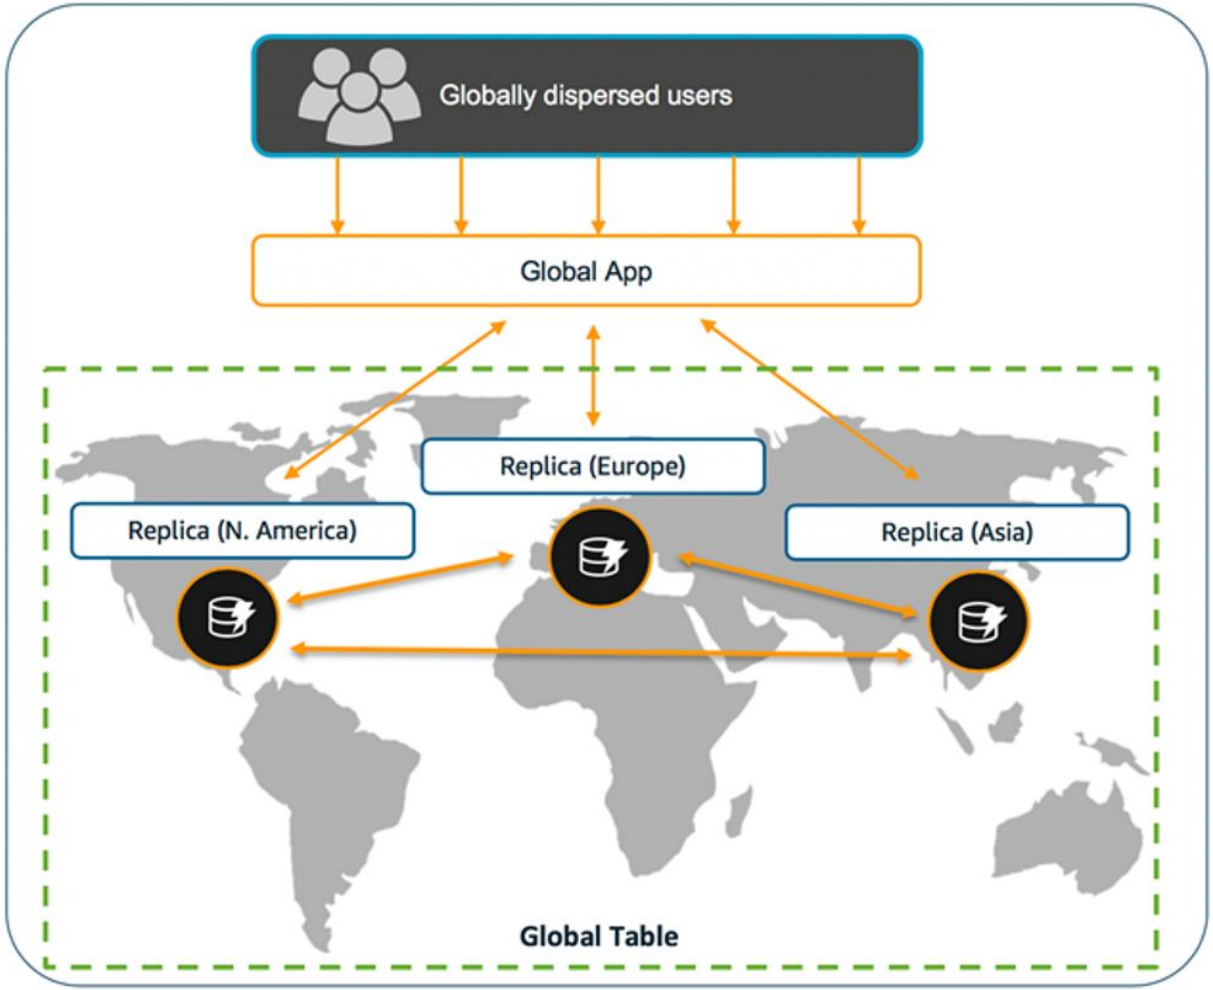
Globally dispersed users
Global App
Replica (Europe)
Replica (N. America)
Replica (Asia)
Global Table

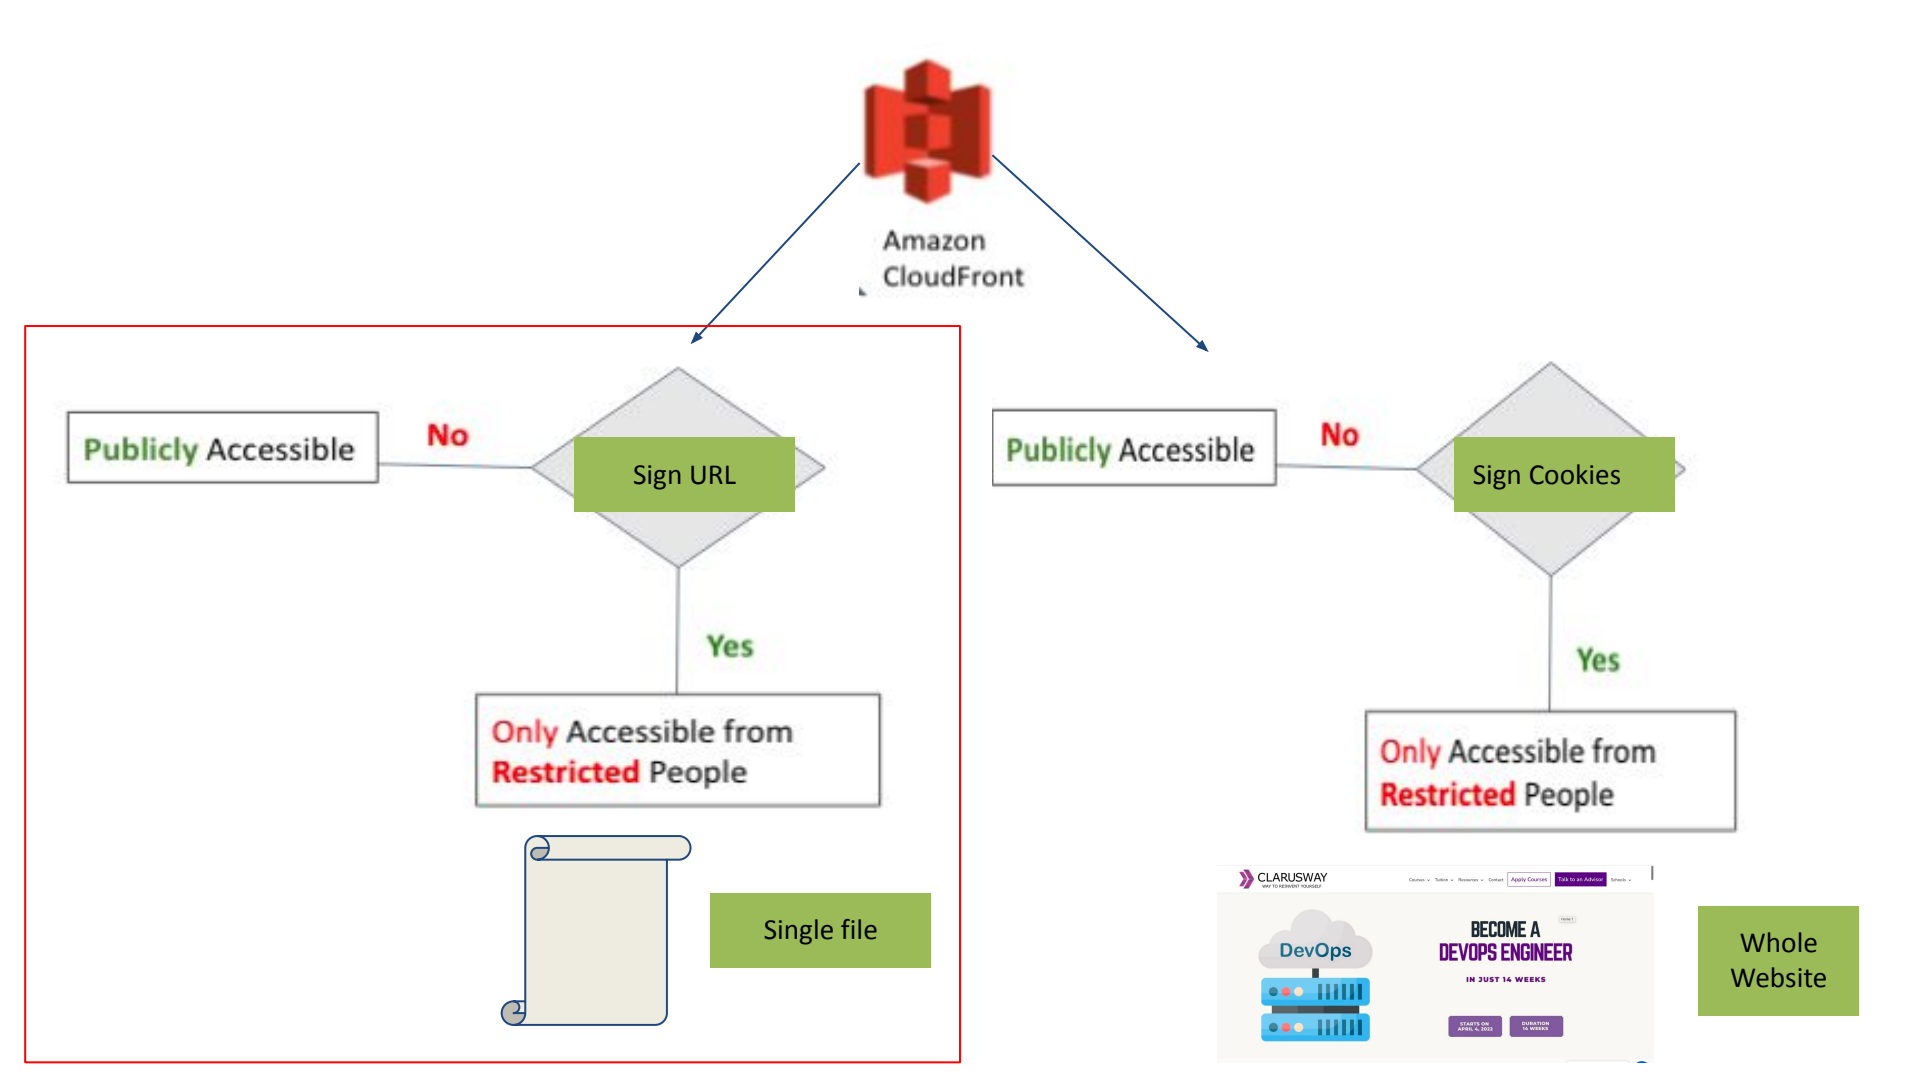
Amazon
CloudFront
Publicly Accessible
No
Sign URL
Yes
Only Accessible from
Restricted People
Single file
Publicly Accessible
No
Sign Cookies
Yes
Only Accessible from
Restricted People
CLARUSWAY
DevOps
BECOME A
DEVOPS ENGINEER
IN JUST 14 WEEKS
Whole
Website

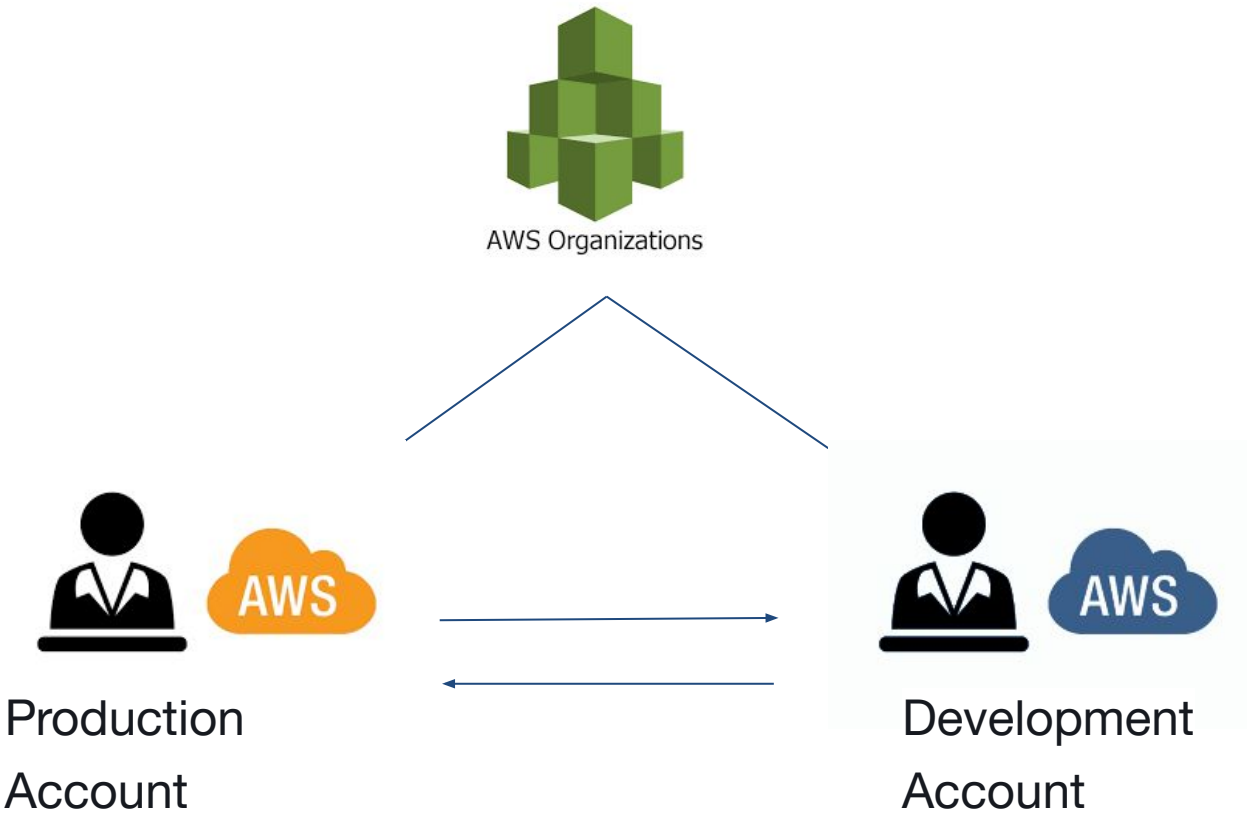

Enable resource sharing in AWS Organizations with RAM

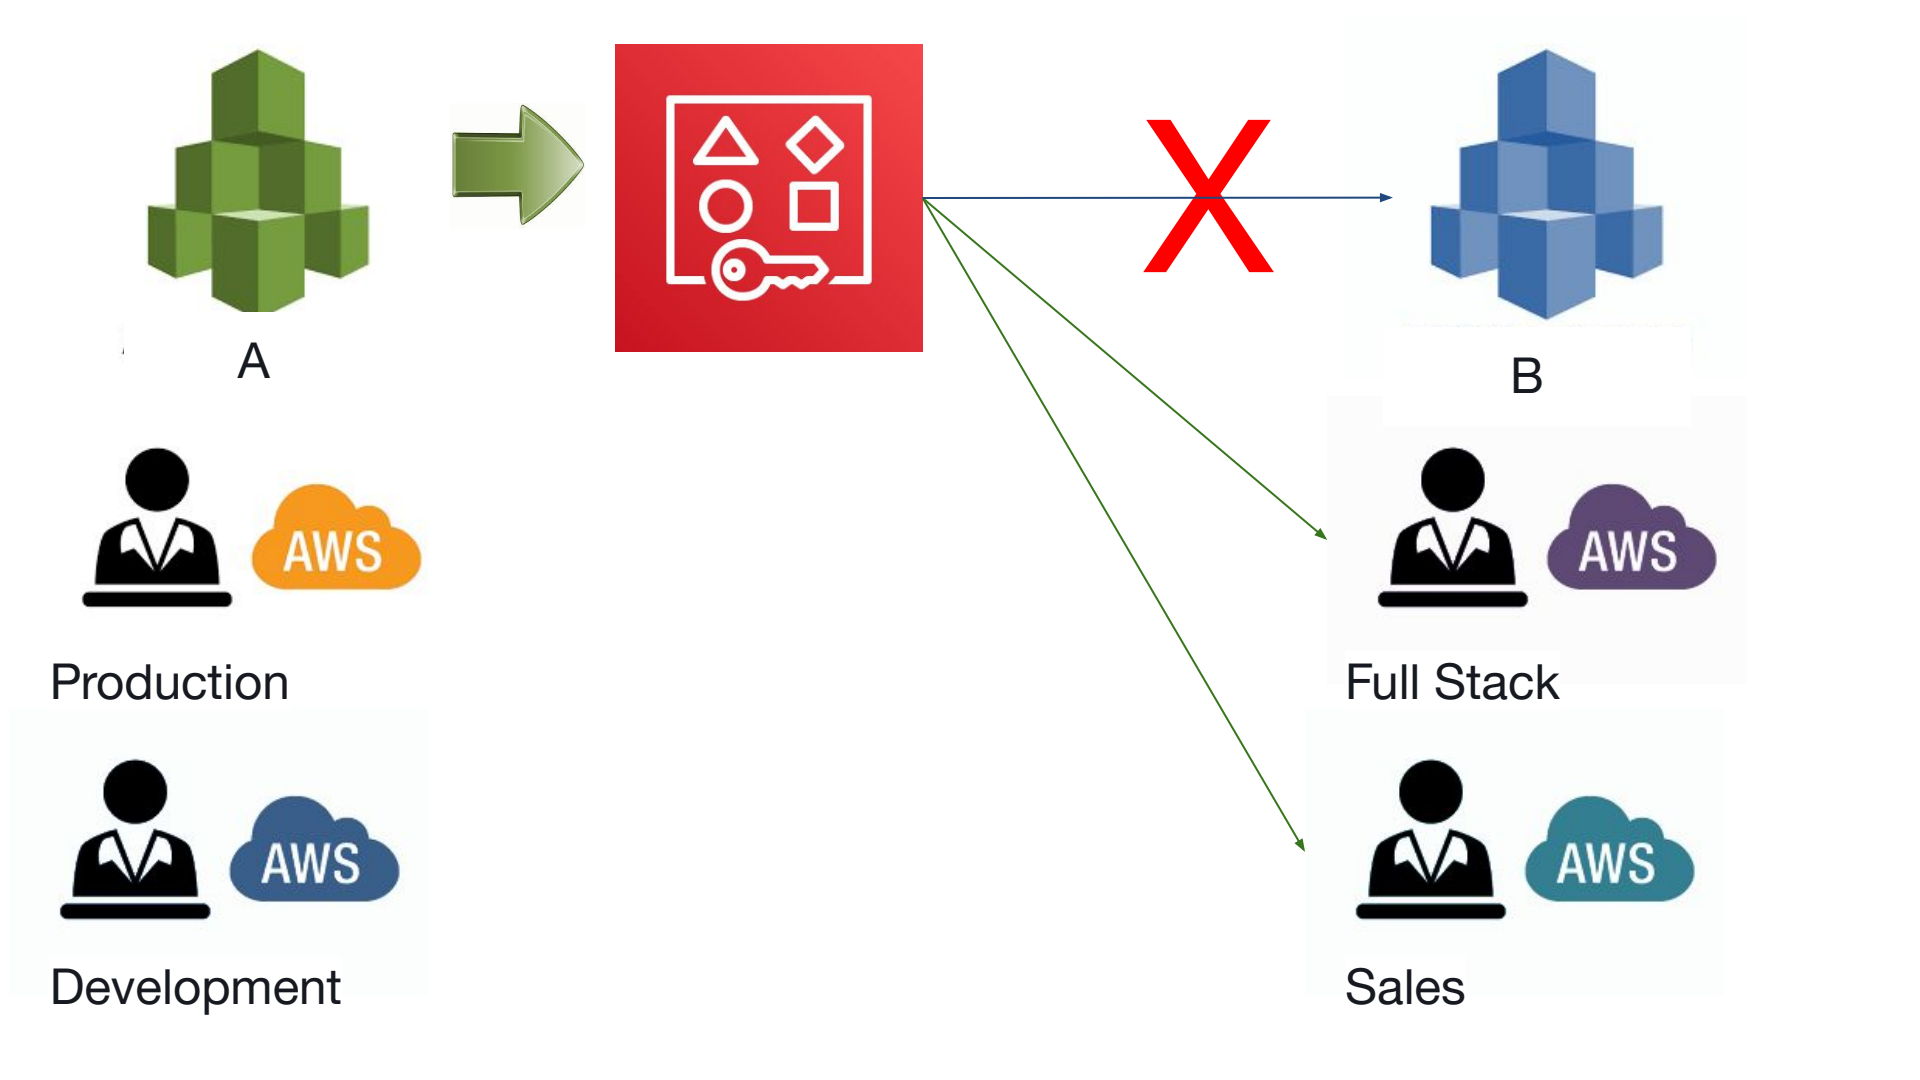
A
B
AWS
Production
AWS
Full Stack
AWS
Development
AWS
Sales

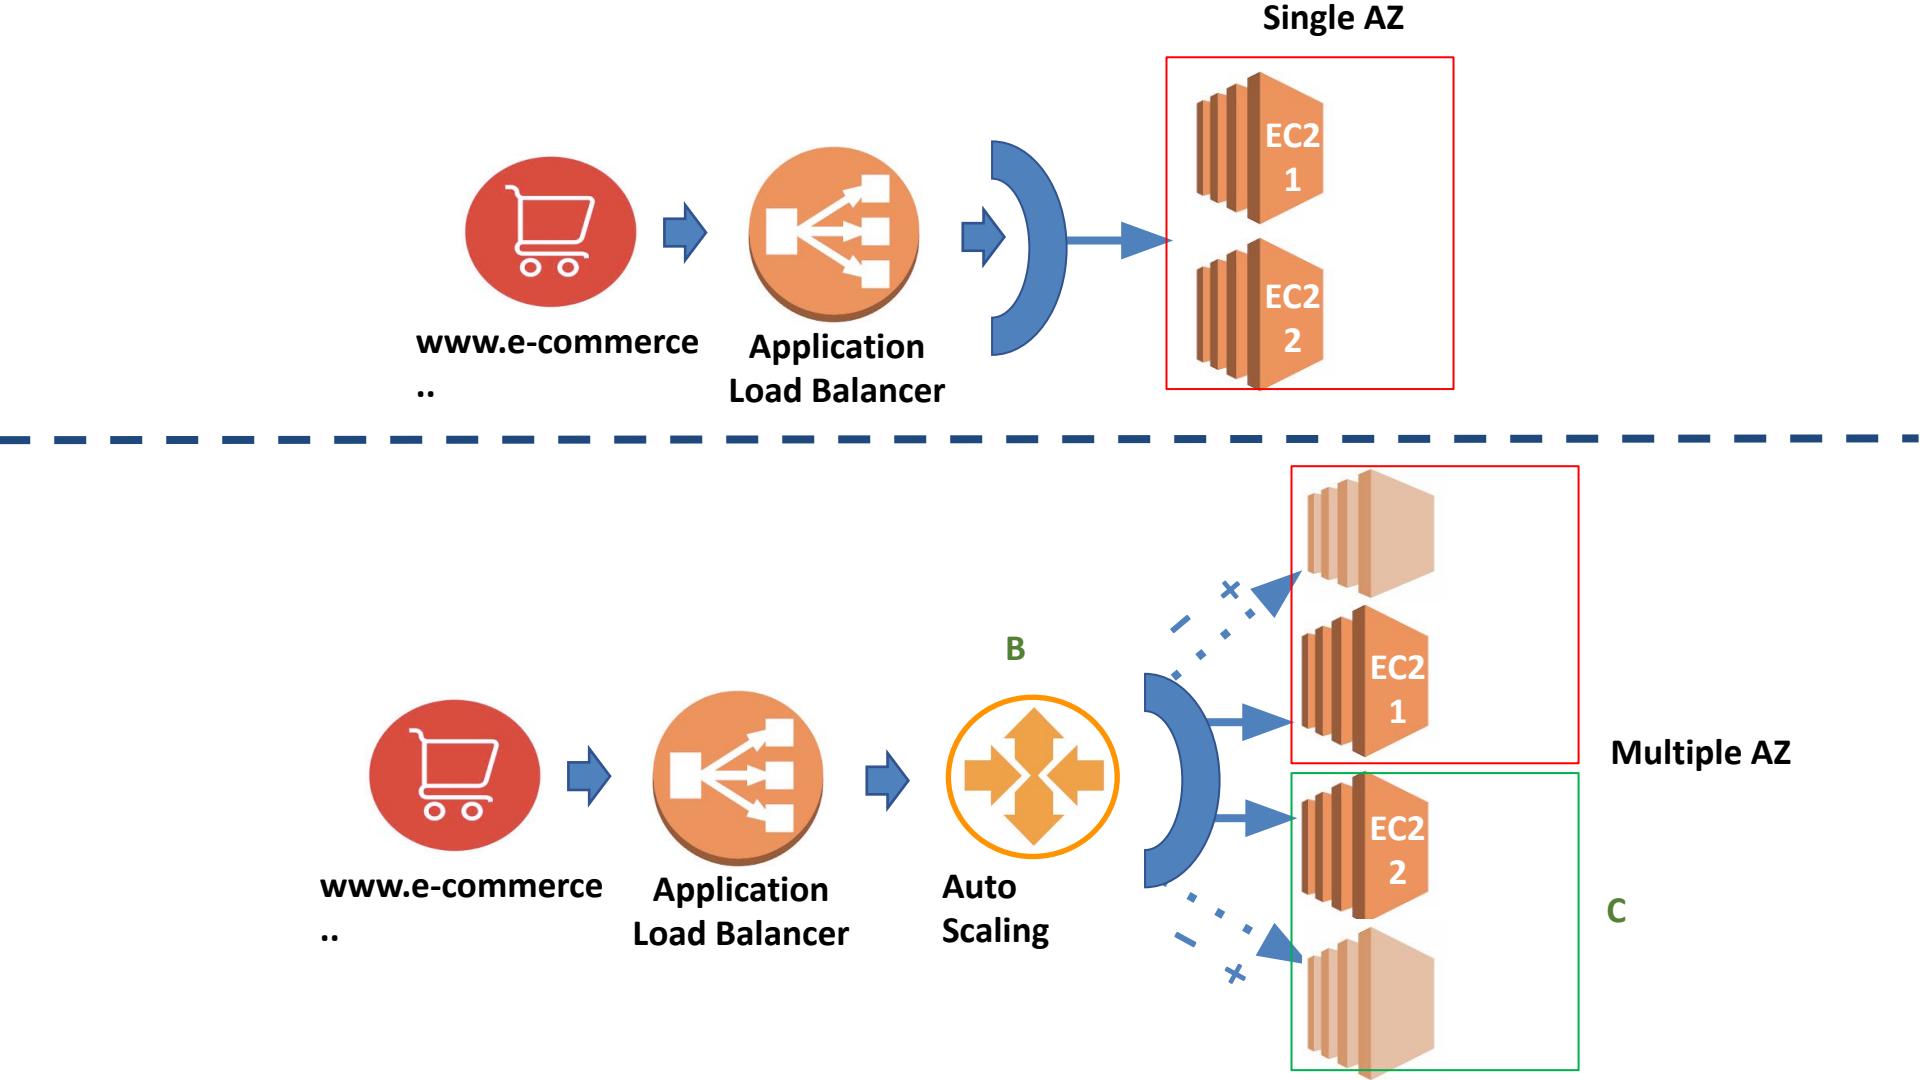
Single AZ
EC2
1
EC2
2
www.e-commerce
..
Application
Load Balancer
B
Auto
Scaling
Multiple AZ
C

backed up your
on-premises data

Storage Gateway

Network File System (NFS)
Server Message Block (SMB).

ISCSI
(Internet Small Computer System

ISCSI

File Gateway

Volume Gateway

Tape Gateway

Stored

Cached

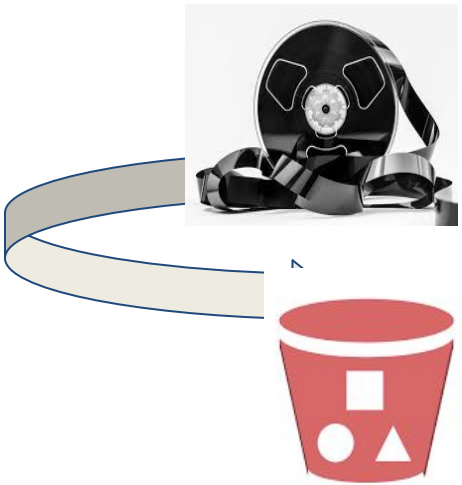

Info: The term **legacy** database commonly refers to a database that has been in use for many years and is therefore unsuitable for modern apps and environments.

- DMS create a replication instance for migrating. Your databases are fully operational when migrating.

- Homogeneous database migrations **(Oracle to Oracle)** **(need Engine conversion)**
Heterogenous database migrations **(Microsoft SQL to Aurora) (Use SCT)**

- Use **SCT (Schema conversion tool)** from one schema to another schema for **Heterogenous migrations**

From Legacy to DynamoDb >>>>>>> Heterogenous database migrations

SCT (Schema conversion tool)

KEY words: Transfer db to AWS ,
Snapshots of db in S3,

No impact on client usage of app
S3 data is fully synced with the on premise db.

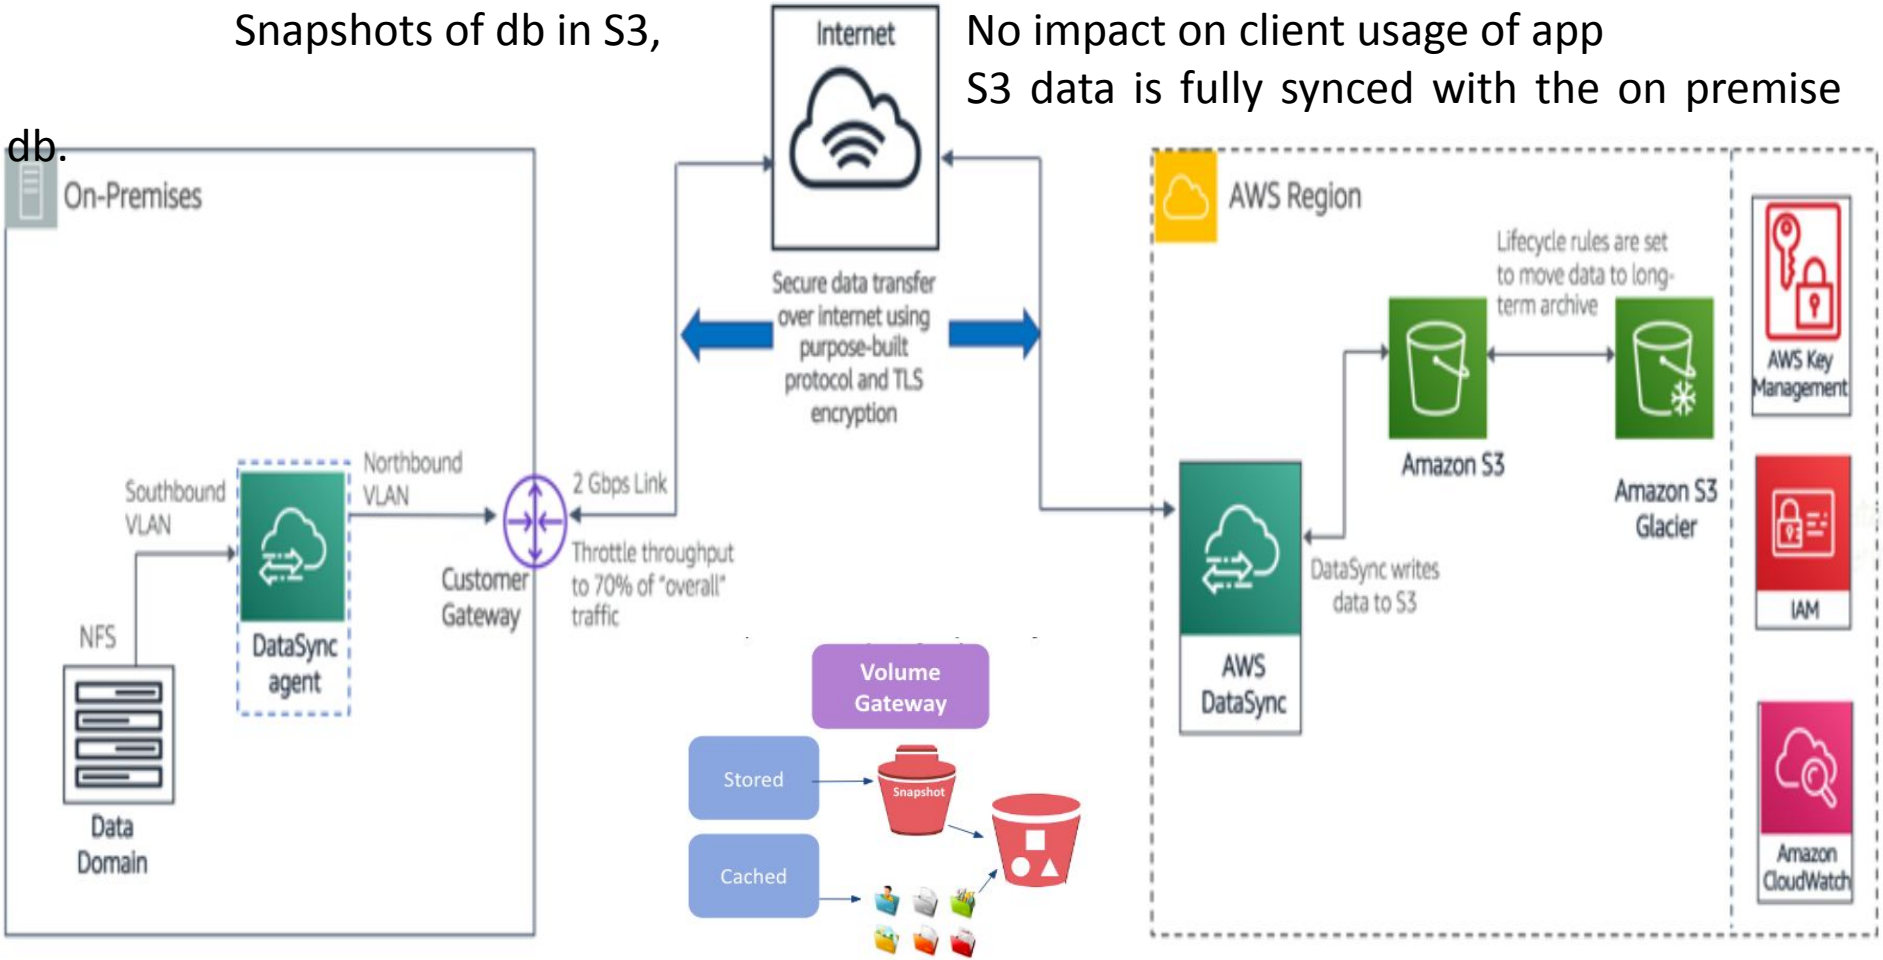

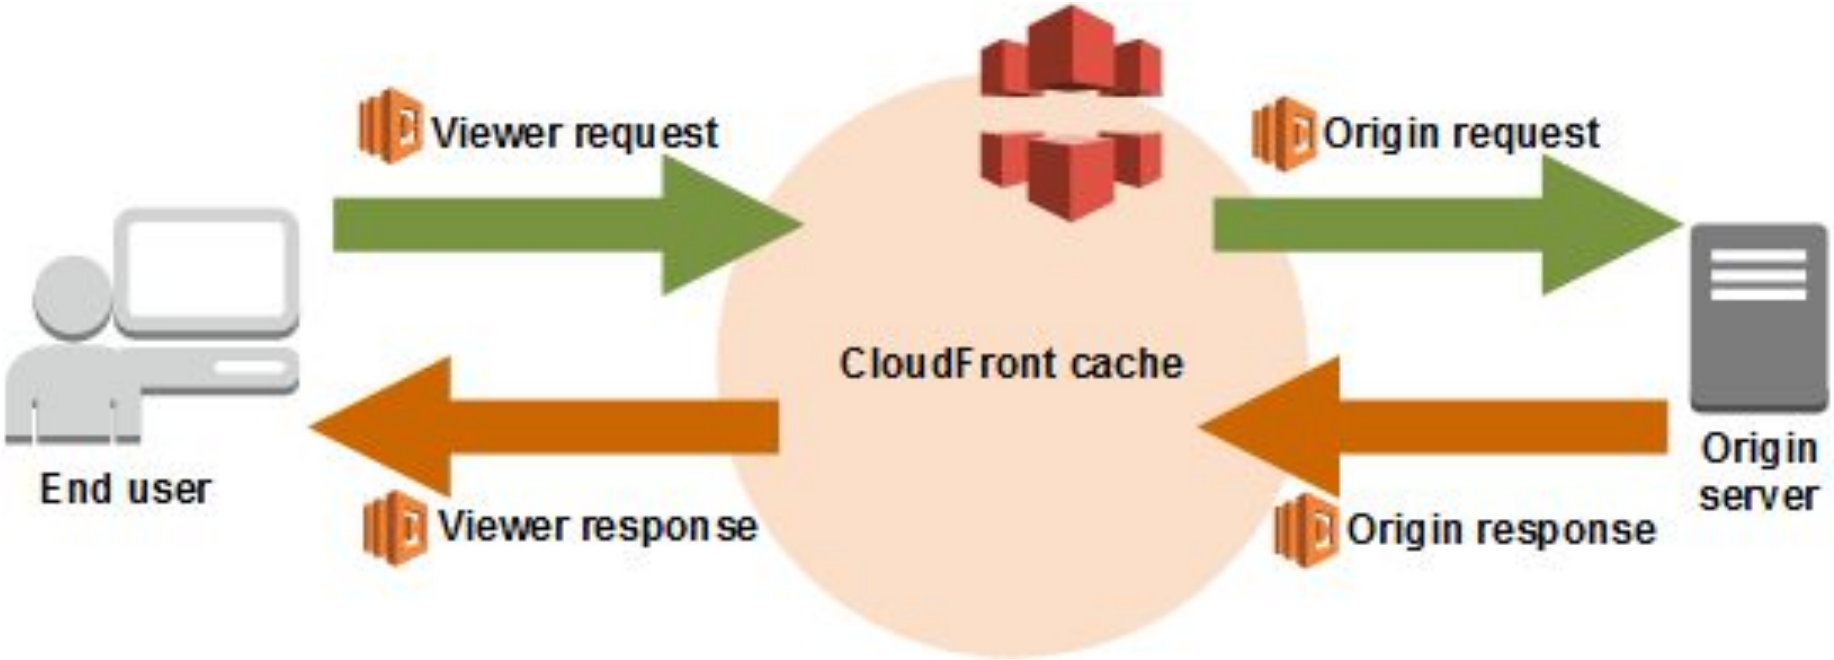
Viewer request
Origin request
CloudFront cache
End user
Origin server
Viewer response
Origin response

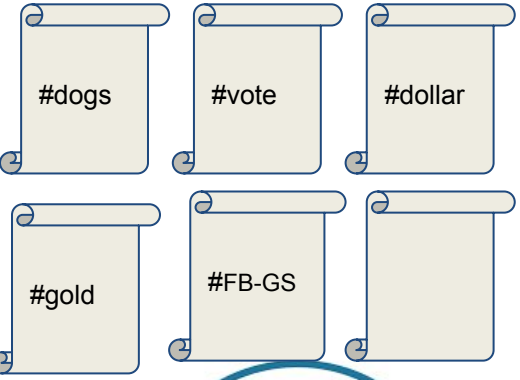
#dogs
#vote
#dollar
#gold
#FB-GS

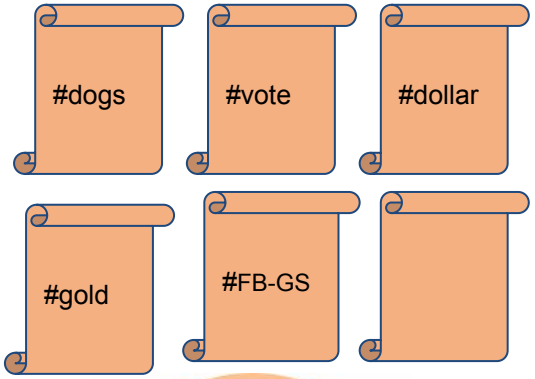
#dogs
#vote
#dollar
#gold
#FB-GS


AWS DataSync


NFS


Amazon EFS

Architecting for the Cloud AWS Best Practices

A stateless application is an application that does not need knowledge of previous interactions and does not store session information. For example, an application that, given the same input, provides the same response to any end user, is a stateless application.

Couple :

Lambda=Stateless

- **Source of the distribution** >>>> OAI >>> from Where
- **Restricted user** >>>> Pre Sign url/cookies >>>> Who (Slide-10)

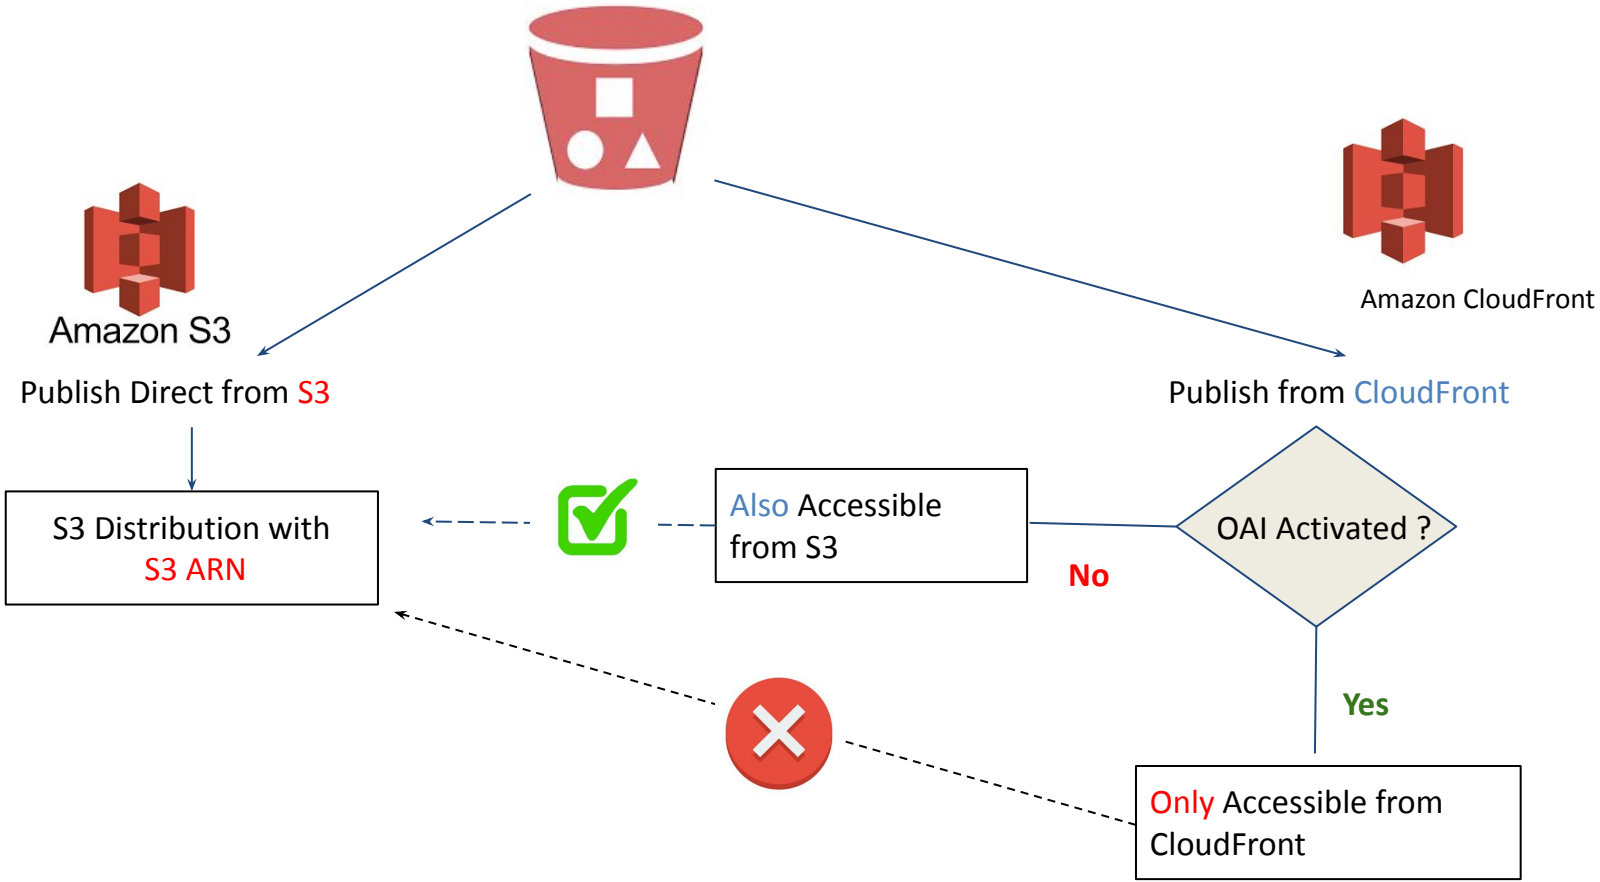

Edit Behavior

Default Cache Behavior Settings

Path Pattern Default (*)

Origin or Origin Group S3-media-sharing-website-bucket

Viewer Protocol Policy
- (•) HTTP and HTTPS
- ( ) Redirect HTTP to HTTPS
- ( ) HTTPS Only

by default it is http and https.

Allowed HTTP Methods
- (•) GET, HEAD
- ( ) GET, HEAD, OPTIONS
- ( ) GET, HEAD, OPTIONS, PUT, POST, PATCH, DELETE

Field-level Encryption Config

Cached HTTP Methods GET, HEAD (Cached by default)

Cache and origin request settings
- (•) Use a cache policy and origin request policy
- ( ) Use legacy cache settings

Cache Policy Managed-CachingOptimized — Create a new policy

View policy details

Learn More

Origin Request Policy — Create a new policy

View policy details

Learn More

Smooth Streaming
- ( ) Yes
- (•) No

Restrict Viewer Access (Use Signed URLs or Signed Cookies)
- ( ) Yes
- (•) No

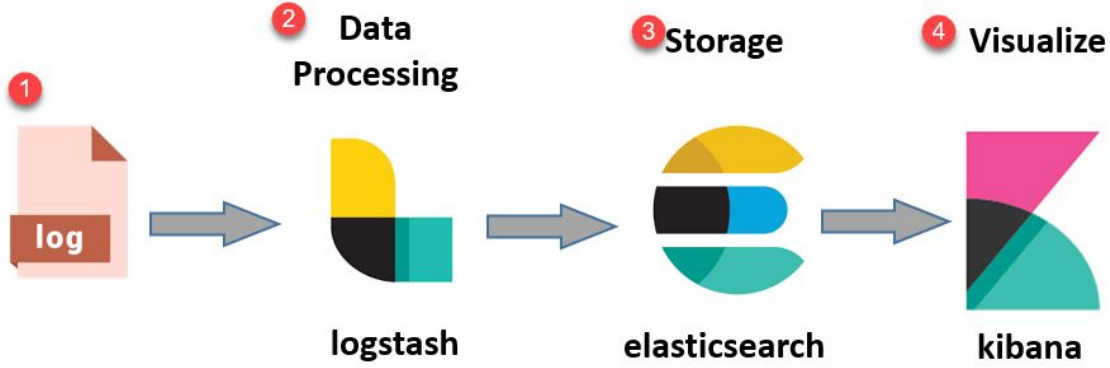

= ELK

Not open source till January 21, 2021

Elasticsearch Service

amazon webservices

AWS Lambda

OpenSearch

Amazon OpenSearch Service

Data Firehouse	- (already built-in support no need lambda)
S3	+ need lambda to be trigger
Cloudwatch Logs	- (already built-in support no need lambda)
AIOT	- (use MQTT)

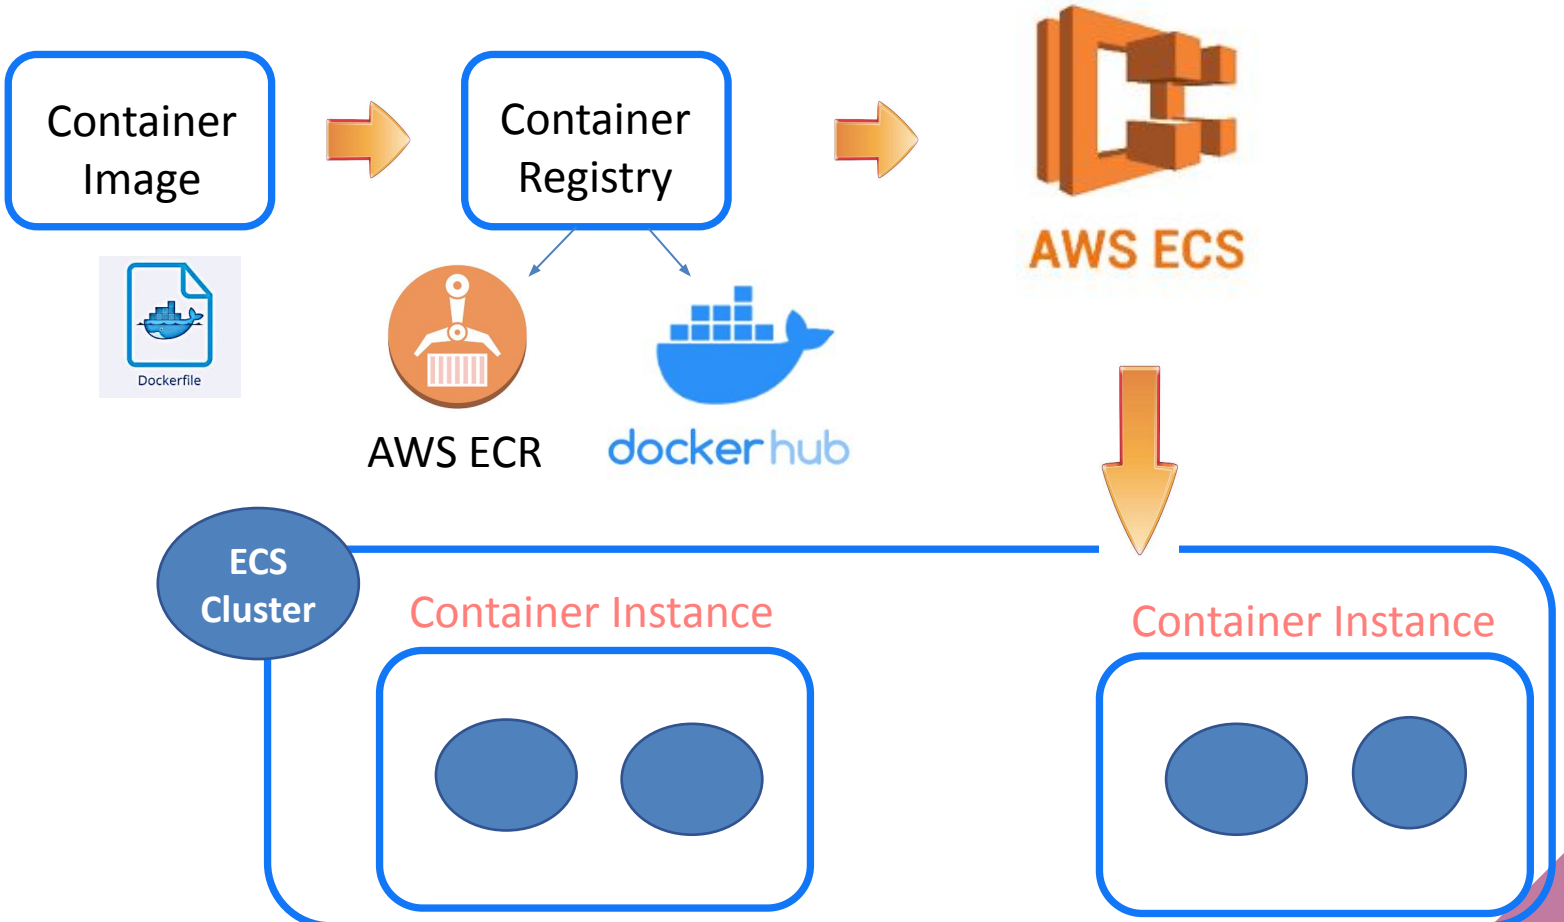
Container Image
Container Registry
AWS ECS
Dockerfile
AWS ECR
dockerhub
ECS Cluster
Container Instance
Container Instance

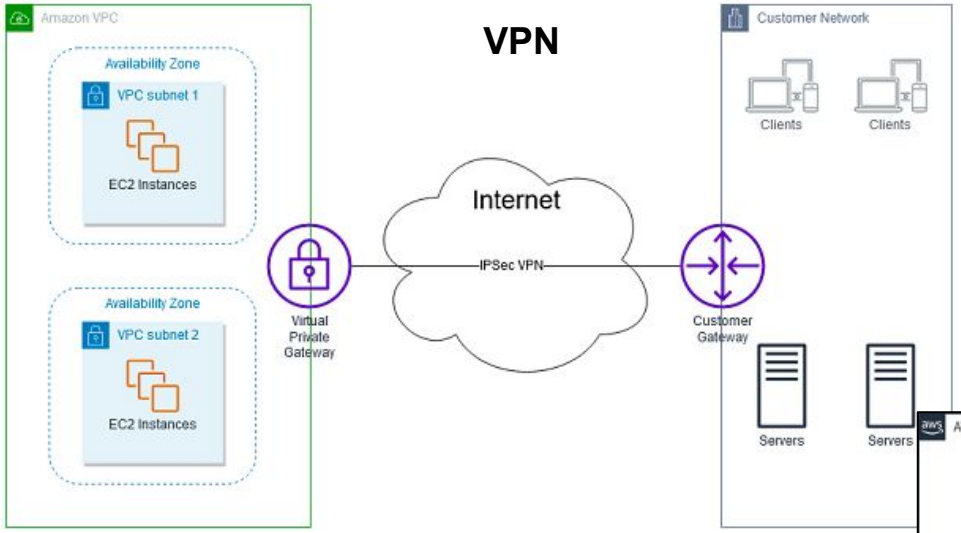

Within the term "IPsec," "IP" stands for "Internet Protocol" and "sec" for "secure."

VPN connections are very cheap (**$37.20**/month
Built in a few minutes

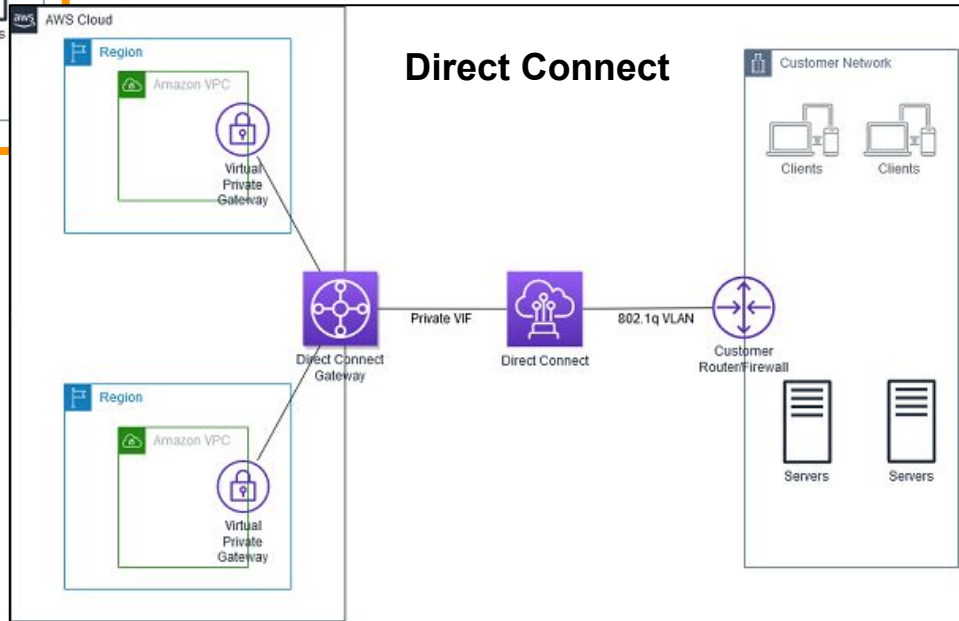

Rotating AWS KMS keys

Automatic key rotation is supported only on symmetric KMS keys

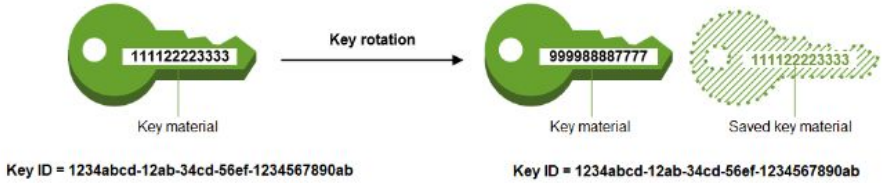

Automatic key rotation is *not* supported on the following types of KMS keys, but you can rotate these KMS keys manually.

- Asymmetric KMS keys
- KMS keys in custom key stores
- KMS keys with imported key material

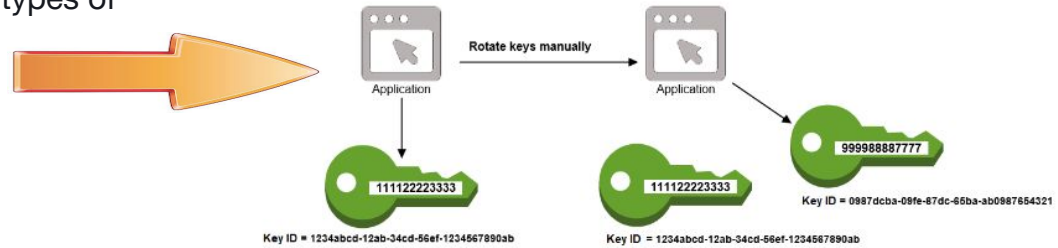

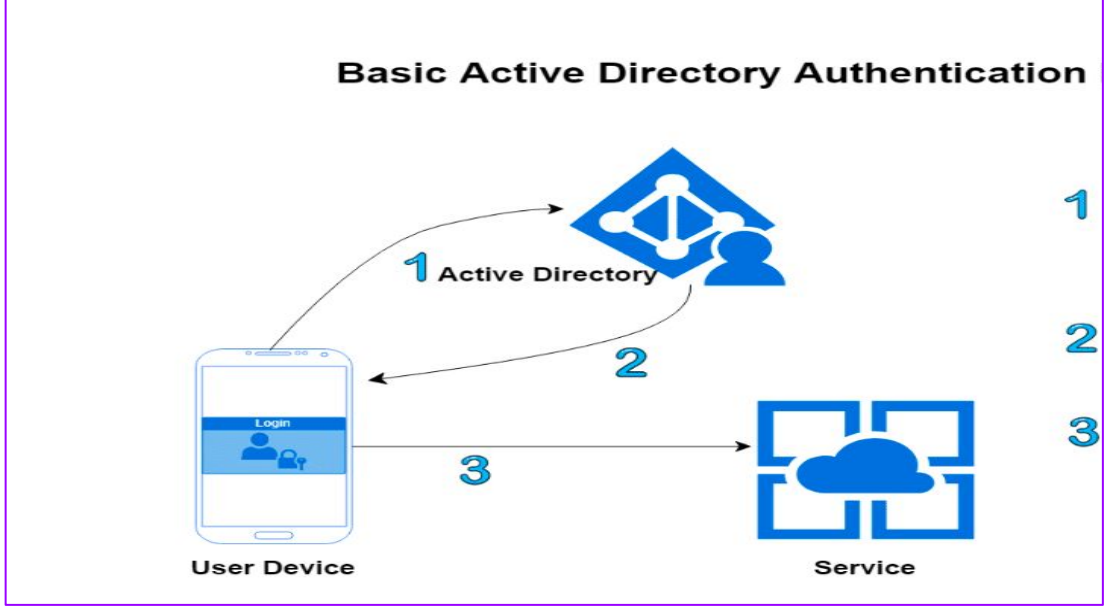

AWS Directory Service for Microsoft Active Directory

Simple AD

AD Connector

Amazon Cognito

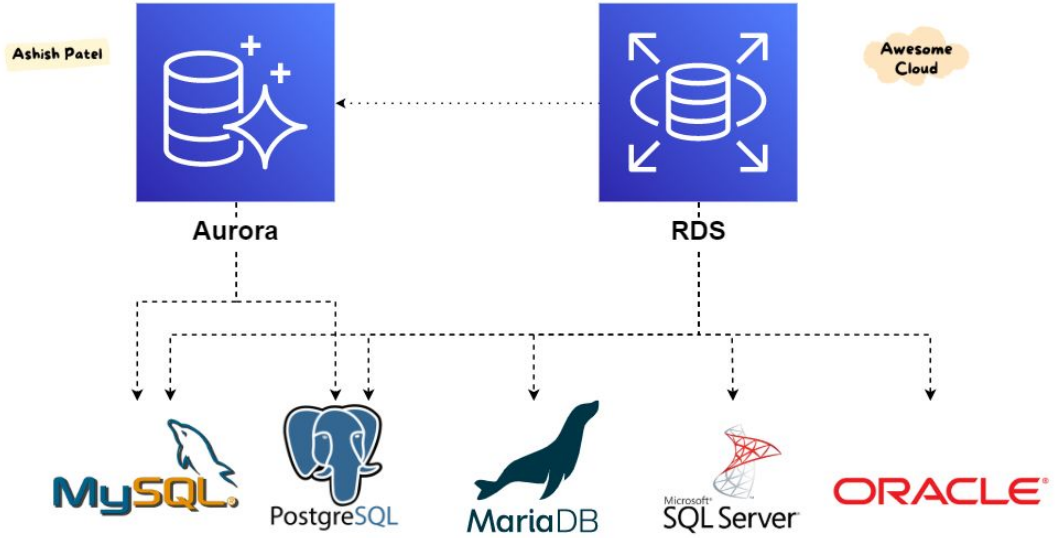

RDS Mysql can grow to 16 Tb
Aurora (Mysql) automatically up to 128 Tb.

In the question : Migrate 5 tb Mysql to AWS. 5Tb will increase in the future.

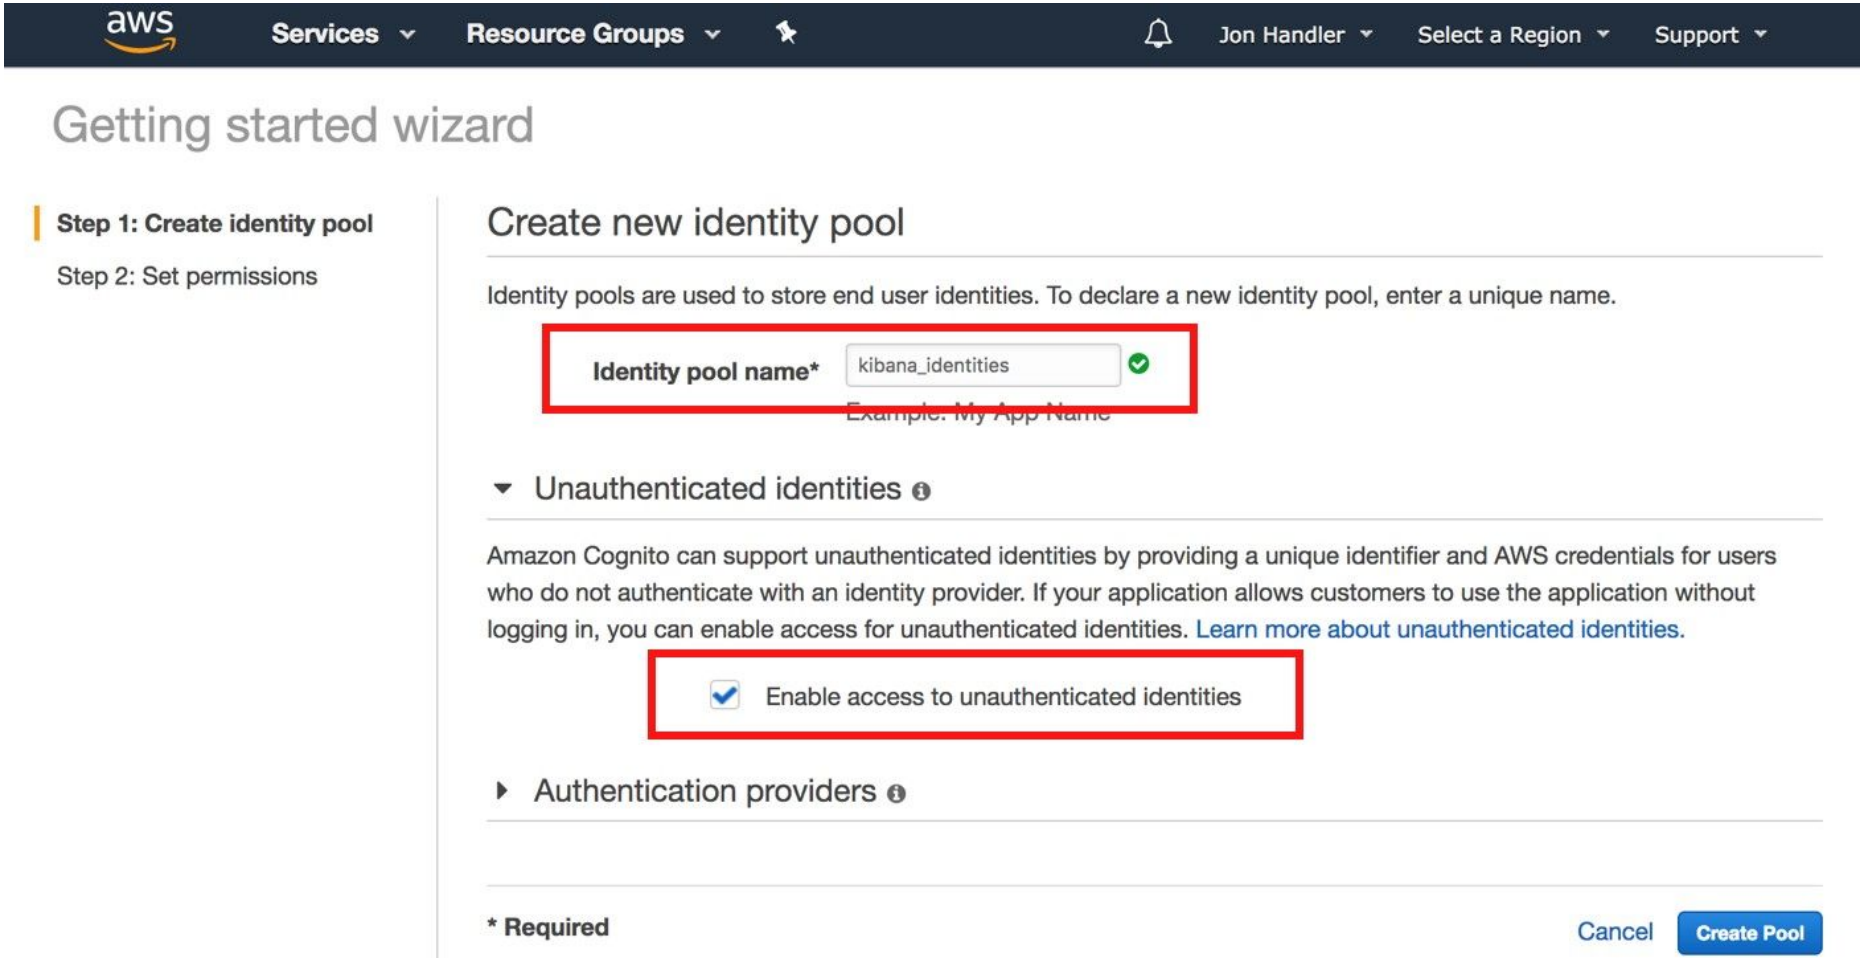
aws
Services
Resource Groups
Jon Handler
Select a Region
Support
Getting started wizard
Step 1: Create identity pool
Step 2: Set permissions
Create new identity pool
Identity pools are used to store end user identities. To declare a new identity pool, enter a unique name.
Identity pool name*
kibana_identities
Example: My App Name
Unauthenticated identities
Amazon Cognito can support unauthenticated identities by providing a unique identifier and AWS credentials for users who do not authenticate with an identity provider. If your application allows customers to use the application without logging in, you can enable access for unauthenticated identities. Learn more about unauthenticated identities.
Enable access to unauthenticated identities
Authentication providers
* Required
Cancel
Create Pool

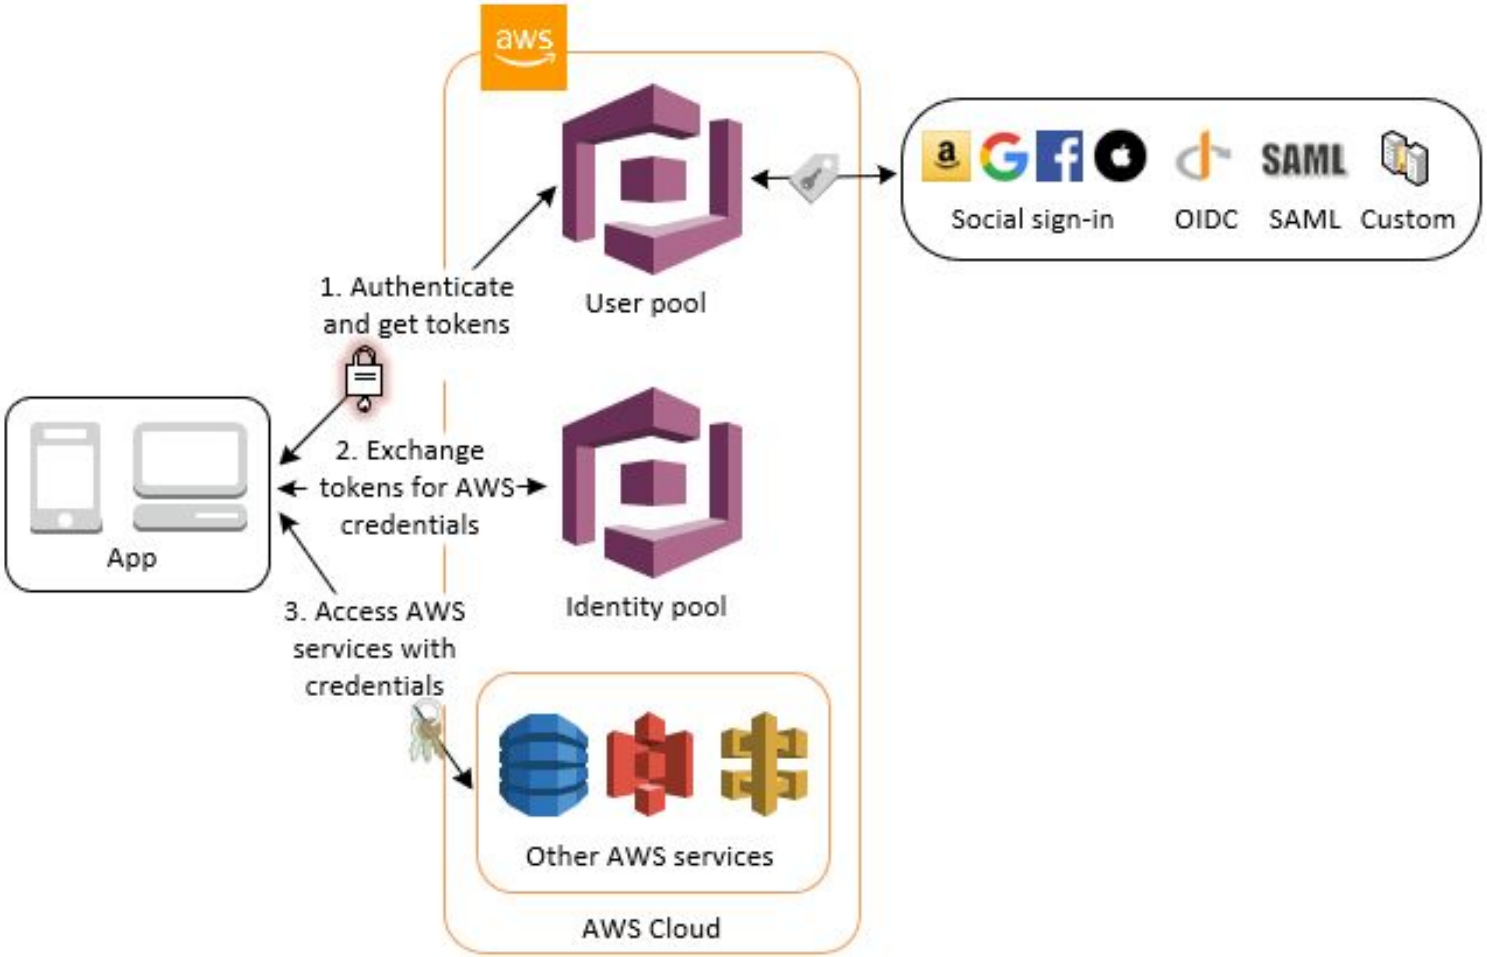
aws
User pool
Social sign-in
OIDC
SAML
Custom
1. Authenticate and get tokens
2. Exchange tokens for AWS credentials
App
Identity pool
3. Access AWS services with credentials
Other AWS services
AWS Cloud

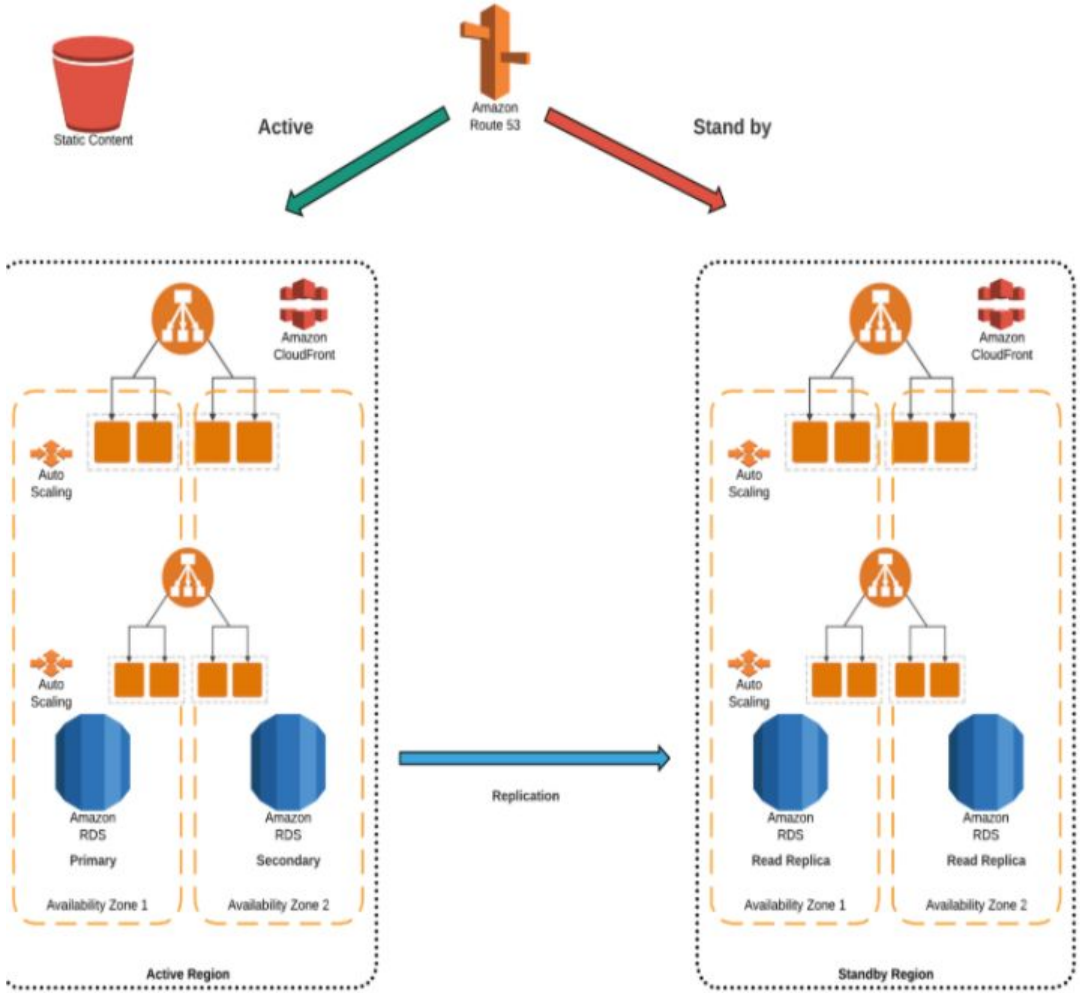

99.95 % SLA

ymore information on the Elastic Load Balancer please refer to the URL below:

User IP: 7.8.9.10/32

Connection Request	
No	Type-Port
1	SSH-22

Security Group-Inbound

Type	Protocol	Port Range	Source
SSH-22	TCP(6)	22	0.0.0.0/0

Inbound Network ACL

Rule	Type	Protocol	Port Range	Source	Allow/ Deny
100	SSH-22	TCP(6)	22	0.0.0.0/0	DENY
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY

Outbound Network ACL

Rule	Type	Protocol	Port Range	Source	Allow/ Deny
100	SSH-22	TCP(6)	22	0.0.0.0/0	ALLOW
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY

User IP: 7.8.9.10/32

Connection Request	
No	Type-Port
1	SSH-22

Security Group-Inbound

Type	Protocol	Port Range	Source
SSH-22	TCP(6)	22	0.0.0.0/0

Inbound Network ACL

Rule	Type	Protocol	Port Range	Source	Allow/ Deny
100	SSH-22	TCP(6)	22	0.0.0.0/0	ALLOW
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY

Outbound Network ACL

Rule	Type	Protocol	Port Range	Source	Allow/ Deny
100	SSH-22	TCP(6)	22	0.0.0.0/0	ALLOW
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY

The 12 requirements of PCI are:

Install and maintain a firewall configuration to protect cardholder data
Do not use vendor-supplied defaults for system passwords and other security parameters
Protect stored cardholder data
Encrypt transmission of cardholder data across open, public networks
Use and regularly update anti-virus software or programs
Develop and maintain secure systems and applications
Restrict access to cardholder data by business need to know
Assign a unique ID to each person with computer access
Restrict physical access to cardholder data
Track and monitor all access to network resources and cardholder data
Regularly test security systems and processes
Maintain a policy that addresses information security for all personnel

PCI Payment Card Industry

CSAA Practice Test 2

Osvaldo

18 Agu 2022